



SIAM EX14 Workshop  
July 7, Chicago - IL

# Preliminary Investigations on Resilient Parallel Numerical Linear Algebra Solvers

Luc GIRAUD

joint work with

E. AGULLO, P. SALAS, E. F. YETKIN, M. ZOUNON

funded by ANR RESCUE and G8-ECS

HiePACS Inria Project  
Joint Inria-CERFACS lab  
INRIA Bordeaux Sud-Ouest

# Context

- ▶ HPC systems are not fault-free
- ▶ A faulty components (node, core, memory) loses all its data
- ▶ Simulations at exascale have to be resilient

Resilience: Ability to compute a correct output in presence of faults

- ▶ Context: Numerical linear algebra
- ▶ Goal: Keep converging in presence of fault
- ▶ Method: Recover-restart strategy without Checkpoint

# Outline

Faults in HPC Systems

Sparse linear systems

Interpolation methods

Numerical experiments

Resilience in eigensolvers

Concluding remarks and perspectives

# Outline

## Faults in HPC Systems

Sparse linear systems

Interpolation methods

Numerical experiments

Resilience in eigensolvers

Concluding remarks and perspectives

# Framework

## Forecast for extreme scale systems

- ▶ Mean Time Between Failure (MTBF): less than one hour
- ▶ Checkpoint time might be larger than MTBF

# Framework

## Forecast for extreme scale systems

- ▶ Mean Time Between Failure (MTBF): less than one hour
- ▶ Checkpoint time might be larger than MTBF

## Objectives

- ▶ Explore fault-tolerant schemes with less/no overhead
- ▶ Numerical algorithms to deal with overhead issue

# Framework

## Forecast for extreme scale systems

- ▶ Mean Time Between Failure (MTBF): less than one hour
- ▶ Checkpoint time might be larger than MTBF

## Objectives

- ▶ Explore fault-tolerant schemes with less/no overhead
- ▶ Numerical algorithms to deal with overhead issue

## Faults in this presentation

- ▶ Detected corrupted memory space (node crashes, damaged memory pages, uncorrected bit-flip, . . . )

# Outline

Faults in HPC Systems

**Sparse linear systems**


Interpolation methods

Numerical experiments

Resilience in eigensolvers

Concluding remarks and perspectives





$$Ax = b$$

We attempt to design fault tolerant solver  
for sparse linear system

## Two classes of iterative methods

- ▶ Stationary methods (Jacobi, Gauss-Seidel, ...)
- ▶ Krylov subspace methods (CG, GMRES, Bi-CGStab, ...)
- ▶ Krylov methods have attractive potential for Extreme-scale

# Outline

Faults in HPC Systems

Sparse linear systems


Interpolation methods

Numerical experiments

Resilience in eigensolvers

Concluding remarks and perspectives


## Block row distribution



We distinguish two categories of data:

- ▶ Static data
- ▶ Dynamic data


## Block row distribution



We distinguish two categories of data:


- ▶ Static data
- ▶ Dynamic data

## Block row distribution




We distinguish two categories of data:

- ▶ Static data
- ▶ Dynamic data



We distinguish two categories of data:

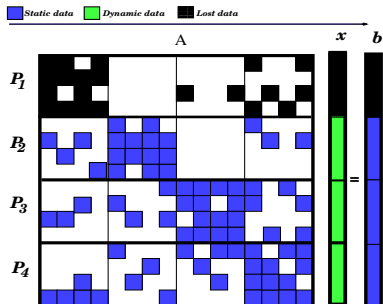
- ▶ Static data
- ▶ Dynamic data



We distinguish two categories of data:

- ▶ Static data
- ▶ Dynamic data

Let's assume that  $P_1$  fails




We distinguish two categories of data:

- ▶ Static data
- ▶ Dynamic data

Let's assume that  $P_1$  fails






We distinguish two categories of data:

- ▶ Static data
- ▶ Dynamic data

Let's assume that  $P_1$  fails

- ▶ Failed processor is replaced
- ▶ Static data are restored




We distinguish two categories of data:

- ▶ Static data
- ▶ Dynamic data

Let's assume that  $P_1$  fails

- ▶ Failed processor is replaced
- ▶ Static data are restored

Reset: Set  $(x_1)$  to initial value



We distinguish two categories of data:


- ▶ Static data
- ▶ Dynamic data

Let's assume that  $P_1$  fails

- ▶ Failed processor is replaced
- ▶ Static data are restored


Our algorithms aim at recovering  $x_1$  and restart

## Overview of our fault tolerant algorithm




- ▶ Sequential simulations
- ▶ Simulation of parallel environment

## Overview of our fault tolerant algorithm




- ▶ Sequential simulations
- ▶ Simulation of parallel environment
- ▶ Generation of fault trace
- ▶ Realistic probability distribution

## Overview of our fault tolerant algorithm




- ▶ Sequential simulations
- ▶ Simulation of parallel environment
- ▶ Generation of fault trace
- ▶ Realistic probability distribution

## Overview of our fault tolerant algorithm



- ▶ Sequential simulations
- ▶ Simulation of parallel environment
- ▶ Generation of fault trace
- ▶ Realistic probability distribution

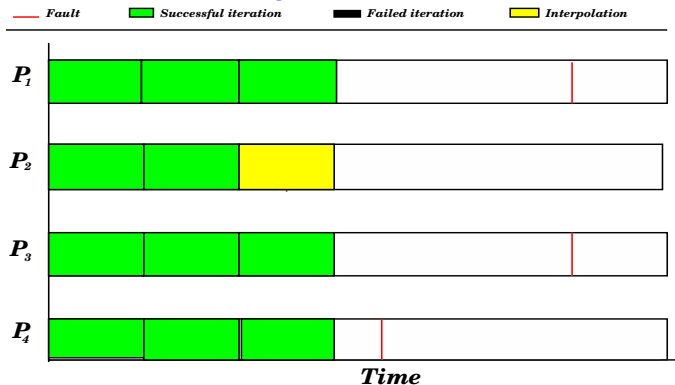
# Overview of our fault tolerant algorithm



- ▶ Sequential simulations
- ▶ Simulation of parallel environment
- ▶ Generation of fault trace
- ▶ Realistic probability distribution




# Overview of our fault tolerant algorithm




- ▶ Sequential simulations
- ▶ Simulation of parallel environment
- ▶ Generation of fault trace
- ▶ Realistic probability distribution

# Overview of our fault tolerant algorithm




- ▶ Sequential simulations
- ▶ Simulation of parallel environment
- ▶ Generation of fault trace
- ▶ Realistic probability distribution

# Overview of our fault tolerant algorithm



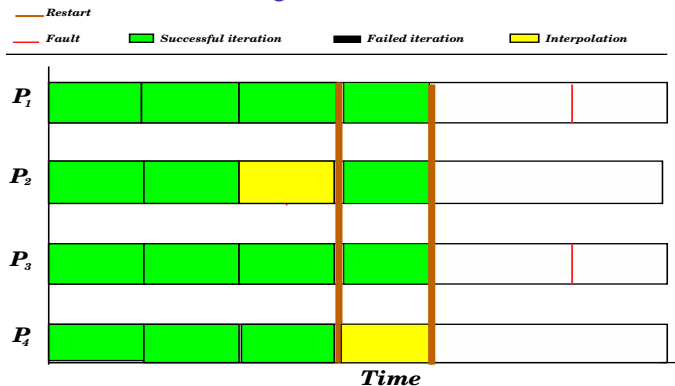
- ▶ Sequential simulations
- ▶ Simulation of parallel environment
- ▶ Generation of fault trace
- ▶ Realistic probability distribution

# Overview of our fault tolerant algorithm



- ▶ Sequential simulations
- ▶ Simulation of parallel environment
- ▶ Generation of fault trace
- ▶ Realistic probability distribution


# Overview of our fault tolerant algorithm



- ▶ Sequential simulations
- ▶ Simulation of parallel environment

- ▶ Generation of fault trace
- ▶ Realistic probability distribution

# Overview of our fault tolerant algorithm



- ▶ Sequential simulations
- ▶ Simulation of parallel environment

- ▶ Generation of fault trace
- ▶ Realistic probability distribution

# Interpolation methods

Fault in linear system

$$\begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \end{pmatrix}$$

# Interpolation methods

Fault in linear system

$$\begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \begin{pmatrix} ? \\ x_2 \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \end{pmatrix} \quad \text{How to recover } x_1?$$



# Interpolation methods

Fault in linear system

$$\begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \begin{pmatrix} ? \\ x_2 \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \end{pmatrix} \quad \text{How to recover } x_1?$$

Linear Interpolation (LI) [Langou, Chen, Bosilca, Dongarra, SISC, 2007]

Solve  $A_{11}x_1 = b_1 - A_{12}x_2$

# Interpolation methods

Fault in linear system

$$\begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \begin{pmatrix} ? \\ x_2 \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \end{pmatrix} \quad \text{How to recover } x_1?$$

Linear Interpolation (LI) [Langou, Chen, Bosilca, Dongarra, SISC, 2007]

$$\text{Solve } A_{11}x_1 = b_1 - A_{12}x_2$$

Least Squares Interpolation (LSI)

$$\begin{pmatrix} A_{11} \\ A_{21} \end{pmatrix} x_1 + \begin{pmatrix} A_{12} \\ A_{22} \end{pmatrix} x_2 = \begin{pmatrix} b_1 \\ b_2 \end{pmatrix}$$

$$x_1 = \underset{x}{\operatorname{argmin}} \left\| \begin{pmatrix} b_1 \\ b_2 \end{pmatrix} - \begin{pmatrix} A_{11} \\ A_{21} \end{pmatrix} x - \begin{pmatrix} A_{12} \\ A_{22} \end{pmatrix} x_2 \right\|_2$$

# Main properties - basic linear algebra

## Proposition

The initial guess generated by LI after a fault does ensure that the A-norm of the forward error associated with the iterates computed by restarted CG or PCG is monotonically decreasing

# Main properties - basic linear algebra

## Proposition

The initial guess generated by LI after a fault does ensure that the A-norm of the forward error associated with the iterates computed by restarted CG or PCG is monotonically decreasing

**[LI might not be defined for non-SPD matrices as diagonal blocks might be singular]**

# Main properties - basic linear algebra

## Proposition

The initial guess generated by LI after a fault does ensure that the A-norm of the forward error associated with the iterates computed by restarted CG or PCG is monotonically decreasing

**[LI might not be defined for non-SPD matrices as diagonal blocks might be singular]**

## Proposition

The initial guess generated by LSI after a fault does ensure the monotonic decrease of the residual norm of minimal residual Krylov subspace methods such as GMRES and MinRES after a restarting due to a failure

# Outline

Faults in HPC Systems

Sparse linear systems

Interpolation methods

**Numerical experiments**

Resilience in eigensolvers

Concluding remarks and perspectives

# Impact of fault rate

Preconditioned GMRES (Kim1 - 2 % data lost)




Figure: 4 faults

# Impact of fault rate

Preconditioned GMRES (Kim1 - 2 % data lost)




Figure: 8 faults



# Impact of fault rate

Preconditioned GMRES (Kim1 - 2 % data lost)




Figure: 17 faults

# Impact of fault rate

Preconditioned GMRES (Kim1 - 2 % data lost)




Figure: 40 faults

# Impact of lost data volume

Preconditioned GMRES(100) (Averous/epb3 - 10 faults)




Figure: 3 % data lost

# Impact of lost data volume

Preconditioned GMRES(100) (Averous/epb3 - 10 faults)




Figure: 0.8 % data lost

# Impact of lost data volume

Preconditioned GMRES(100) (Averous/epb3 - 10 faults)




Figure: 0.2 % data lost

# Impact of lost data volume

Preconditioned GMRES(100) (Averous/epb3 - 10 faults)




Figure: 0.001 % data lost

# Penalty of restart strategy

- ▶ Recover-restart strategy
- ▶ When restarting, we lose the Krylov subspace built before the fault
- ▶ Consequence: delay of convergence due to restart
- ▶ Restarting mechanism is naturally implemented in GMRES to reduce the computational resource consumption
- ▶ CG does not need to be restarted

# Penalty of restart strategy on PCG




Figure: PCG on a 7-point stencil 3D Poisson equation with 70 faults - 5 % data lost



# Penalty of restart strategy on PCG




Figure: PCG on a 7-point stencil 3D Poisson equation with 70 faults - 5 % data lost

# Outline

Faults in HPC Systems

Sparse linear systems

Interpolation methods

Numerical experiments

Resilience in eigensolvers

Concluding remarks and perspectives

# Recovery-restart for eigensolvers

Fault in eigenproblem

$$\begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \lambda \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$$

# Recovery-restart for eigensolvers

Fault in eigenproblem

$$\begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \begin{pmatrix} ? \\ x_2 \end{pmatrix} = \lambda \begin{pmatrix} ? \\ x_2 \end{pmatrix} \quad \text{How to recover } x_1?$$


Linear Interpolation (LI)

Solve the linear system  $(A_{11} - \lambda I_1) x_1 = -A_{12}x_2$

Least Squares Interpolation (LSI)

$$\begin{pmatrix} A_{11} \\ A_{21} \end{pmatrix} x_1 + \begin{pmatrix} A_{21} \\ A_{22} \end{pmatrix} x_2 = \lambda \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$$

$$x_1 = \underset{x}{\operatorname{argmin}} \left\| \begin{pmatrix} A_{11} - \lambda I_1 \\ A_{21} \end{pmatrix} x + \begin{pmatrix} A_{12} \\ A_{22} - \lambda I_2 \end{pmatrix} x_2 \right\|_2$$




If  $Ax = \lambda x$  with  $x \neq 0$ , where  $A \in \mathbb{C}^{n \times n}$ ,  $x \in \mathbb{C}^n$ , and  $\lambda \in \mathbb{C}$ , then,

- ▶  $\lambda$  : eigenvalue
- ▶  $x$  : eigenvector
- ▶  $(\lambda, x)$  : eigenpair

## Two classes of methods

- ▶ Fixed Point Methods (Power Method, Subspace iteration)
- ▶ Subspace Methods (Jacobi-Davidson, Arnoldi, IRA/Krylov Schur)

# Thermo-acoustic test example



# Jacobi-Davidson method




Figure: Jacobi-Davidson method with 5 faults - 1 % lost data. Convergence history using LSI and Checkpoint of current iterate

# Outline

Faults in HPC Systems

Sparse linear systems

Interpolation methods

Numerical experiments

Resilience in eigensolvers

Concluding remarks and perspectives



# Concluding remarks

## Summary

- ▶ We have designed techniques to interpolate meaningful lost data based on simple linear algebra tools
- ▶ Our techniques preserve some of the key monotonicity of Krylov solvers but lack of robustness of LI for non-SPD problems
- ▶ The restarting effect remains reasonable within the GMRES context
- ▶ No fault, no overhead
- ▶ These techniques can be adapted to multiple faults
- ▶ What about silent soft-error - CGPOP preliminary experiments ?

Merci for your attention

Questions ?



<https://team.inria.fr/hiepacs/>