

Technical Report about implementation in CUDA of Monte Carlo Linear Solvers (MCLS)

Steven Hamilton, Massimiliano Lupo Pasini

Abstract

In this report timings are presented as concerns Monte Carlo linear Solvers (MCLS) with different configurations. Different test cases are used to get quite a general overview of the behavior. In particular various options to generate random numbers as well as use of texture memory are tested for the sake of efficiency. All the test cases have been run three times to filter the fluctuations due to hardware issues.

1 Different random number generation techniques

The generation of random numbers plays an important role in the calculation of the solution to a linear system through a Monte Carlo procedure, since a new random number must be generated to either step ahead in a random walk or to kink off a new one. Since this operation must be repeated very frequently in the code, it is recommendable to decrease its cost as much as possible, not to affect the efficiency of the overall performance. Two viable options to initialize random number generators are:

1. Having the same seed but different sequence number generates a number guaranteed to be 2^{67} away from each other, but the downside is the heavy computations to advance the 2^{67} position

```
__global__ void initialize_rng(curandState *state, int seed, int offset)
{
    int tid = threadIdx.x + blockIdx.x * blockDim.x;

    curand_init(seed, tid, offset, &state[tid]);
}
```

2. Giving different seeds, and just keep the sequence number at 0, it's a lot faster but there might be correlation between threads, since there is no guarantee on the separation between each threads

```
__global__ void initialize_rng2(curandState *state, int*seed, int offset)
{
    int tid = threadIdx.x + blockIdx.x * blockDim.x;
```

```

    curand_init(seed[tid], 0, offset, &state[tid]);
}

```

Matrix	Size	Nb. Histories	% time rng	% time kernel	tot. time (ms)	rel. residual norm
2D Laplacian	900	10^7	36.0	64.0	299,563	0.0931241
1D Laplacian	10^6	10^4	0.05	99.5	50,016	12.3509
1D Laplacian	10^6	10^7	18.3	81.7	640,563	0.391686
SP ₁	25,568	10^7	18.0	82.0	349,096	0.450792

Table 1: Timings using same seed and different sequences.

Matrix	Size	Nb. Histories	% time rng	% time kernel	tot. time (ms)	rel. residual norm
2D Laplacian	900	10^7	0.0	100.0	191,796	0.119588
1D Laplacian	10^6	10^4	0.0	100.0	50,094	12.3914
1D Laplacian	10^6	10^7	0.0	100.0	533,451	0.408134
SP ₁	25,568	10^7	0.0	100.0	241,571	0.397887

Table 2: Timings using different seeds and same sequence.

As it can be noticed from the previous tables 1 and 2, the second option provides better timings. However the payoff is an increase of the final relative residual, likely due to the correlation between the sequences generated by different seeds. The quantities "time rng" and "time kernel" are expressed in terms of percentage of the total time.

2 Use of texture memory

The LDG instruction (exposed via the `__ldg` intrinsic) is a memory load that uses the texture path. It has the advantage that it does not require the explicit use of textures, since it does not explicitly bind one. Therefore `__ldg()` reads data through the texture path, without requiring a texture itself. It is an overloaded function with the prototype `__ldg(const *T)` where T is one of CUDA's built-in types. The perk of using LDG instruction is that explicit uses of textures causes a certain amount of code clutter and overhead (e.g. for API calls to bind textures). Classical textures also use the texture load path, but in addition can transform both index (e.g. clamping modes) and data returned (e.g. interpolation) in various ways; the necessary control information is provided to the hardware during texture binding. Because the texture cache is non-coherent with respect to writes in the same kernel, use of the texture load path requires that the underlying data is read-only across the entire kernel.

3 LDG calls into the code

LDG instructions have been introduced in the following functions of Profugus code:

1. `lower_bound`
2. `initialize_history`
3. `getNewState`

4. tallyContribution

All the instructions employed above provide a runtime decrease. However, some of them are more significant than the others. In particular the most effective calls are the ones located in 1) and 3). This is due to the fact that these functions are the ones called most frequently. The option of using LDG instruction at point 4) actually plays a role just when the expected value estimator is employed.

In the examples represented below the length of the history is set to 10000 steps. The simulations have been accomplished both with and without a weight cut off equal to 10^{-9} .

Matrix	Size	Nb. Histories	$\rho(H)$	$\rho(\hat{H})$	tot. time (ms)
2D Laplacian	900	10^7	0.994869	0.99447	191,818
1D Laplacian	10^6	10^4	0.799972	0.639977	35,351
1D Laplacian	10^6	10^7	0.799972	0.639977	517,788
SP ₁	25,568	10^7	0.977674	0.999836	349,096

Table 3: Timings without LDG instructions and constant history length equal to 10,000.

Matrix	Size	Nb. Histories	$\rho(H)$	$\rho(\hat{H})$	tot. time (ms)
2D Laplacian	900	10^7	0.994869	0.99447	191,818
1D Laplacian	10^6	10^4	0.799972	0.639977	35,351
1D Laplacian	10^6	10^7	0.799972	0.639977	517,788
SP ₁	25,568	10^7	0.977674	0.999836	349,096

Table 4: Timings without LDG instructions and constant history length equal to 10,000.

Matrix	Size	Nb. Histories	$\rho(H)$	$\rho(\hat{H})$	tot. time (ms)
2D Laplacian	900	10^7	0.994869	0.99447	139,882
1D Laplacian	10^6	10^4	0.799972	0.639977	37,610
1D Laplacian	10^6	10^7	0.799972	0.639977	500,092
SP ₁	25,568	10^7	0.977674	0.999836	157,545

Table 5: Timings with LDG instructions at 1) and constant history length equal to 10,000.

Matrix	Size	Nb. Histories	$\rho(H)$	$\rho(\hat{H})$	tot. time (ms)
2D Laplacian	900	10^7	0.994869	0.99447	193,336
1D Laplacian	10^6	10^4	0.799972	0.639977	37,759
1D Laplacian	10^6	10^7	0.799972	0.639977	548,909
SP ₁	25,568	10^7	0.977674	0.999836	241,490

Table 6: Timings with LDG instructions at 2) and constant history length equal to 10,000.

By comparing the values of tables 3-13 it is pointed out that a benefit in terms of timings is achieved overall when the `ldg` instruction is employed in all the subroutines. Moreover the improvements are more evident for the 2D laplacian and the SP₁ matrix. The reason of this might be associated with the sparsity pattern. Indeed a higher number of nonzero entries induces the

Matrix	Size	Nb. Histories	$\rho(H)$	$\rho(\hat{H})$	tot. time (ms)
2D Laplacian	900	10^7	0.994869	0.99447	105,010
1D Laplacian	10^6	10^4	0.799972	0.639977	36,340
1D Laplacian	10^6	10^7	0.799972	0.639977	519,408
SP ₁	25,568	10^7	0.977674	0.999836	200,559

Table 7: Timings with LDG instructions at 3) and constant history length equal to 10,000.

Matrix	Size	Nb. Histories	$\rho(H)$	$\rho(\hat{H})$	tot. time (ms)
2D Laplacian	900	10^7	0.994869	0.99447	140,347
1D Laplacian	10^6	10^4	0.799972	0.639977	36,344
1D Laplacian	10^6	10^7	0.799972	0.639977	499,007
SP ₁	25,568	10^7	0.977674	0.999836	157,458

Table 8: Timings with LDG instructions at 1), 2) and constant history length equal to 10,000.

Matrix	Size	Nb. Histories	$\rho(H)$	$\rho(\hat{H})$	tot. time (ms)
2D Laplacian	900	10^7	0.994869	0.99447	158,468
1D Laplacian	10^6	10^4	0.799972	0.639977	36,141
1D Laplacian	10^6	10^7	0.799972	0.639977	547,777
SP ₁	25,568	10^7	0.977674	0.999836	200,325

Table 9: Timings with LDG instructions at 2), 3) and constant history length equal to 10,000.

Matrix	Size	Nb. Histories	$\rho(H)$	$\rho(\hat{H})$	tot. time (ms)
2D Laplacian	900	10^7	0.994869	0.99447	105,331
1D Laplacian	10^6	10^4	0.799972	0.639977	34,616
1D Laplacian	10^6	10^7	0.799972	0.639977	520,762
SP ₁	25,568	10^7	0.977674	0.999836	118,200

Table 10: Timings with LDG instructions at 1), 2), 3) and constant history length equal to 10,000.

Matrix	Size	Nb. Histories	$\rho(H)$	$\rho(\hat{H})$	tot. time (ms)
2D Laplacian	900	10^7	0.994869	0.99447	27,686
1D Laplacian	10^6	10^4	0.799972	0.639977	35,108
1D Laplacian	10^6	10^7	0.799972	0.639977	40,720
SP ₁	25,568	10^7	0.977674	0.999836	4,919

Table 11: Timings without LDG instructions and with weight cutoff.

histories to run for longer. Since the histories keep running for more steps in these cases, the utility of the texture memory might increase as well. For the 2D laplacian and the SP₁ matrix the employment of `ldg` instructions almost halves the time for the computation. The time reduction gets weaker for the other test cases. In fact for the 1D laplacian it seems there is no benefit in terms of timings coming from the texture memory.

By looking at the results of the same test cases when a weight cutoff is used, we can see that the utility of `__ldg` instructions is vanished (see Tables 14-19).

Matrix	Size	Nb. Histories	$\rho(H)$	$\rho(\hat{H})$	tot. time (ms)
2D Laplacian	900	10^7	0.994869	0.99447	26,344
1D Laplacian	10^6	10^4	0.799972	0.639977	35,911
1D Laplacian	10^6	10^7	0.799972	0.639977	41,092
SP ₁	25,568	10^7	0.977674	0.999836	4,759

Table 12: Timings with LDG instructions at 1) and with weight cutoff.

Matrix	Size	Nb. Histories	$\rho(H)$	$\rho(\hat{H})$	tot. time (ms)
2D Laplacian	900	10^7	0.994869	0.99447	27,722
1D Laplacian	10^6	10^4	0.799972	0.639977	36,005
1D Laplacian	10^6	10^7	0.799972	0.639977	40,001
SP ₁	25,568	10^7	0.977674	0.999836	5,585

Table 13: Timings with LDG instructions at 2) and with weight cutoff.

Matrix	Size	Nb. Histories	$\rho(H)$	$\rho(\hat{H})$	tot. time (ms)
2D Laplacian	900	10^7	0.994869	0.99447	24,814
1D Laplacian	10^6	10^4	0.799972	0.639977	36,604
1D Laplacian	10^6	10^7	0.799972	0.639977	40,305
SP ₁	25,568	10^7	0.977674	0.999836	5,065

Table 14: Timings with LDG instructions at 3) and with weight cutoff.

Matrix	Size	Nb. Histories	$\rho(H)$	$\rho(\hat{H})$	tot. time (ms)
2D Laplacian	900	10^7	0.994869	0.99447	26,248
1D Laplacian	10^6	10^4	0.799972	0.639977	35,741
1D Laplacian	10^6	10^7	0.799972	0.639977	40,017
SP ₁	25,568	10^7	0.977674	0.999836	4,701

Table 15: Timings with LDG instructions at 1) and 2) and with weight cutoff.

Matrix	Size	Nb. Histories	$\rho(H)$	$\rho(\hat{H})$	tot. time (ms)
2D Laplacian	900	10^7	0.994869	0.99447	27,184
1D Laplacian	10^6	10^4	0.799972	0.639977	36,141
1D Laplacian	10^6	10^7	0.799972	0.639977	40,727
SP ₁	25,568	10^7	0.977674	0.999836	5,086

Table 16: Timings with LDG instructions at 2), 3) and with weight cutoff.

4 Generation of random numbers

Random numbers must be employed for each random walk to:

- determine what is the initial state
- determine the following state in the path accordingly to the current one

Essentially the generation of random numbers is located in the subroutines:

Matrix	Size	Nb. Histories	$\rho(H)$	$\rho(\hat{H})$	tot. time (ms)
2D Laplacian	900	10^7	0.994869	0.99447	24,307
1D Laplacian	10^6	10^4	0.799972	0.639977	36,980
1D Laplacian	10^6	10^7	0.799972	0.639977	40,596
SP ₁	25,568	10^7	0.977674	0.999836	4,921

Table 17: Timings with LDG instructions at 1), 2) and 3) and with weight cutoff.

1. `initializeHistory`

2. `getNewState`

These two operations have to be repeated for all the random walks employed in the computations. In order to minimize the time spent in the generation of random numbers, it might be useful to gather the generation of many of these at the same time. Therefore a gathering of the random number generator's calls has been accomplished, partially modifying the subroutines "`initializeHistory`" and "`getNewState`". The relationship one-to-one between the generation of a random number and a call to one of these two subroutines is broken. A group of random numbers is generated in advance before the use of the aforementioned functions. The size of the batch of this grouped random numbers is a tuning parameter that can be set up to find the optimal configuration, reducing the overall time of execution.

The viable options that might be adopted are:

1. to call separately the random number generator once for the `initializeHistory` and then employ the batch for the successive steps
2. to start employing the batch even for the initial step of the history

Matrix	Size	Nb. Histories	% time rng	% time kernel	tot. time (ms)	rel. residual norm
2D Laplacian	900	10^7	0.0	100.0	191,796	0.119588
1D Laplacian	10^6	10^4	0.0	100.0	50,094	12.3914
1D Laplacian	10^6	10^7	0.0	100.0	533,451	0.408134
SP ₁	25,568	10^7	0.0	100.0	241,652	0.366509

Table 18: Results for separate calls of rng for `initializeHistory` and `getNewState`.

Matrix	Size	Nb. Histories	% time rng	% time kernel	tot. time (ms)	rel. residual norm
2D Laplacian	900	10^7	0.0	100.0	192,325	0.119588
1D Laplacian	10^6	10^4	0.0	100.0	36,602	12.3914
1D Laplacian	10^6	10^7	0.0	100.0	517,536	0.408134
SP ₁	25,568	10^7	0.0	100.0	241,928	0.397887

Table 19: Results for a single call of rng for both `initializeHistory` and `getNewState`.

Overall, for a configuration where the length of the random walk is fixed, the second option seems to be slightly more efficient (compare Table 18 and 19).