

MONTE CARLO METHODS TO SOLVE SPARSE
LINEAR SYSTEMS

Massimiliano Lupo Pasini

Contents

1	Formulation of plain Monte Carlo linear solvers	3
1.1	Forward Method	4
1.2	Adjoint Method	5
2	Statistical requirements	6
2.1	Forward method - Expected value and variance of the estimator	6
2.2	Adjoint method - Expected value and variance of the estimator	7
2.3	Numerical interpretation of the statistical requirements	7
2.4	Necessary conditions, sufficient conditions and choice of the transition matrix . . .	8
3	Adaptive approaches	10
3.1	An adaptive cut-off for the length of random walks	10
3.1.1	Relative standard deviation - Forward method	10
3.1.2	Relative standard deviation - Adjoint method	11
3.1.3	An a posteriori variance-based adaptivity	11
3.2	Numerical results	11
3.3	A stationarity check for the a posteriori adaptivity	12
3.4	An a priori adaptive approach on the residual	13
4	Linear solvers with hybrid schemes	15
4.1	Monte Carlo Synthetic Acceleration	15
4.2	Preconditioning techniques	15
4.2.1	Block Preconditioning	15
4.2.2	Approximate Inverse of a Sparse Matrix (AINV)	16
4.2.3	Incomplete LU Factorization Preconditioners (ILU)	17
4.3	Reordering techniques	18
4.4	Numerical results	18
4.4.1	Illustrative examples	18
4.4.2	SP_N equations	18
5	Applicable sets for MC solvers	23
5.1	Strictly diagonally dominant matrices	23
5.1.1	Numerical results	24
5.2	Generalized diagonally dominant matrices (GDDM)	25
5.3	"Block-diagonally dominant"-like matrices	26
5.3.1	Forward method	26
5.3.2	Adjoint method	27
	References	27

Introduction

For a long period the pursuit of a reliable solver for linear and nonlinear algebraic systems has been one of the most important issues in Numerical Analysis. Above all for mathematical problems derived from practical issues by modeling. The mathematical properties of the model (e.g. likely sparsity pattern of the matrix) have urged scientists to look for efficient and pattern-driven solvers, capable to cope with curse of dimensionality and high scalability requests.

As scalability computing move towards exascale facilities (which stands for machines computing up to 10^{18} flops), new numerical schemes suitable for this kind of devices are strongly demanded. In fact the purpose is to combine high fidelity of the results with the possibility to exploit the hardware enrichment that computer industry is going to provide us with.

Recently new algorithms that combine statistical and numerical devices have been developed. Even though *numerical efficiency* has been sacrificed sometimes, the advantages in terms of *robustness* of these methods make it worth carrying on studies in this direction.

This report is written in order to show some results concerning the resolution of linear systems via a Monte Carlo approach. The work presented here is based on what developed so far on this topic by some of the authors who in literature gave a contribution (see [Hal62], [Hal94], [DA98], [DAK01], [AAD⁺05], [ESW13] and [EMSH14]).

In Chapters 1 the main issue is a comparison between two different approaches, both aiming at solving an algebraic problem stochastically. These two approaches, one transversal to the other, are named *Forward* and *Adjoint Method*. The basic idea is to generate random walks with random variables associated with them. The role of the random variables is to reconstruct the values attained by the solution vector to the linear problem we are interested in.

In Chapter 2 theoretical constraints are presented. These must be respected for convergence issues, referring to results published in [HMY13]. Sufficient conditions, necessary conditions and equivalence conditions for convergence are stated.

Chapter 3 is dedicated to issues associated with implementation and computation. In particular an adaptive approach for the cut-off of the random walks is presented, based on what is described in [Sla13]. Afterwards an adaptive selection of histories to employ is introduced. Basic numerical tests are employed for illustrative purposes.

Chapter 4 is focused on numerical schemes elaborated through the years in order to speed-up the convergence of plain Monte Carlo linear solvers. Specifically formulations of *Sequential Monte Carlo* and *Monte Carlo Synthetic Acceleration* are treated. Illustrative numerical results and more practical ones are included.

Chapter 5 is targeted to combine all the topics covered in the previous chapters to come up with an overview about the state-of-the-art plight.

Chapter 1

Formulation of plain Monte Carlo linear solvers

As already said, the goal of this family of methods is to use random walks for estimating the solution vector of a linear system by interpreting it as the expected value of a random variable (the estimator).

As underlined in [DA98], Monte Carlo methods may be divided into two classes: *direct methods* ([Hal62], [Hal94], [DA98], [DAK01]) and *iterative* ones ([ESW13] and [EMSH14]). The first are characterized by a merely stochastic scheme, therefore the provided error with respect to the exact solution is made of just a stochastic component. The second ones, instead, combine numerical and statistical schemes and they generate two types of error: one is *stochastic*, the other one is *systematic*. It does not mean that it will be always simple to recognize them separately, but it is important to be aware of this intrinsic structure in order to target what is the part of the scheme that requires a finer refinement (e.g. increase of numerical iterations rather than random walks).

We start our treatise by introducing the direct methods.

Assume to start from a linear system such as

$$A\mathbf{x} = \mathbf{b}, \quad (1.1)$$

where $A \in \mathbb{R}^{n \times n}$, $\mathbf{x}, \mathbf{b} \in \mathbb{R}^n$.

By applying a left preconditioning, the system (1.1) gets the following shape:

$$P^{-1}A\mathbf{x} = P^{-1}\mathbf{b}, \quad (1.2)$$

1.2 can be reinterpreted as a fixed point scheme

$$\mathbf{x} = H\mathbf{x} + \mathbf{f}. \quad (1.3)$$

where $H = I - P^{-1}A$ and $\mathbf{f} = P^{-1}\mathbf{b}$.

Assuming that the spectral radius $\rho(H) < 1$, the solution to (1.3) can be written in terms of a power series of H (Neumann-Ulam series)

$$\mathbf{x} = \sum_{i=0}^{\infty} H^i \mathbf{f}.$$

Therefore the fixed point scheme generates a sequence of approximate solutions $\{\mathbf{x}^{(k)}\}_{k=0}^{\infty}$ which converges to the exact solution regardless of the initial guess \mathbf{x}_0 .

By restricting the attention to a single component of \mathbf{x} we have

$$x_i = \sum_{\ell=0}^{\infty} \sum_{k_1=1}^n \sum_{k_2=1}^n \cdots \sum_{k_{\ell}=1}^n H_{k_0, k_1} H_{k_1, k_2} \cdots H_{k_{\ell-1}, k_{\ell}} f_{k_{\ell}}. \quad (1.4)$$

The last equation can be reinterpreted as the realization of an estimator defined on a random walk. Let us start considering a random walk whose state space S is characterized by the set of indexes of the forcing term \mathbf{f} .

$$S = \{1, 2, \dots, n\} \subset \mathbb{N}.$$

At each i -th step of the random walk it is associated a random variable k_i . Its realization represents the index of the component of \mathbf{f} which is visited in the current step of the random walk.

The way the transition probability is built and the way the initial state of the random walk is chosen give birth to the two different approaches we are going to use.

1.1 Forward Method

The goal is to evaluate a functional such as

$$J(\mathbf{x}) = (\mathbf{h}, \mathbf{x}) = \sum_{i=1}^n h_i x_i.$$

where $\mathbf{h} \in \mathbb{R}^n$ is the Riesz representative in \mathbb{R}^n of the functional J . We can use it to build the initial probability $\tilde{p} : S \rightarrow [0, 1]$ of the random walk such that

$$\tilde{p}(k_0 = i) = \tilde{p}_{k_0} = \frac{|h_i|}{\sum_{i=1}^n |h_i|}.$$

It is important to stress out that the role of vector \mathbf{h} is restricted to the building of the initial probability. Once that is done, it is not used anymore in the definition of the stochastic process. A possible choice for the transition probability P instead can be

$$p(k_i = j \mid k_{i-1} = i) = P_{i,j} = \frac{|H_{i,j}|}{\sum_{k=1}^n |H_{i,k}|}.$$

F where $pr(\cdot, i) : S \rightarrow [0, 1] \forall i \in S$. A related sequence of random variables $w_{i,j}$ can be defined such that

$$w_{i,j} = \frac{H_{i,j}}{P_{i,j}}.$$

The probability distribution of the random variables $w_{i,j}$ is represented by the transition matrix that rules the stochastic process. The sequence of random variables just introduced can be used to build one more sequence of random variables, whose expected value will assume the role of estimator. At first we introduce quantities W_j such that

$$W_0 = \frac{h_{k_0}}{\tilde{p}_{k_0}}, \quad W_j = W_{j-1} w_{i,j}, \quad j = 1, \dots, i.$$

By defining

$$X_i(\nu) = \sum_{m=0}^k W_m f_{i_m}$$

as the random variable associated with a specific permutation ν is the random walk indexed by i , we can introduce the estimator θ_i such as

$$\theta_i = E[X_i] = \sum_{\nu} P_{\nu} X_i(\nu), \quad i = 1, \dots, n.$$

Integer n represents the size of the solution vector to 1.1 and index i is referred to the component of the solution vector we want to compute. P_{ν} is the probability associated with a specific permutation of the random walk. It can be proved that

$$E[W_i f_{k_i}] = (\mathbf{h}, H^i \mathbf{f}), \quad i = 0, 1, 2, \dots$$

and

$$\theta_i = E \left[\sum_{i=0}^{\infty} W_i f_{k_i} \right] = (\mathbf{h}, \mathbf{x}).$$

A possible choice for \mathbf{h} is a vector of the canonical basis. This would correspond in setting manually the initial state of the random walk because the related initial probability is a Dirac delta

$$\tilde{p}(k_0 = i) = \delta_{i,j}.$$

By doing so, we have

$$\theta_i = E \left[\sum_{\ell=0}^{\infty} W_{\ell} f_{k_{\ell}} \right] = x_i = \sum_{l=0}^{\infty} \sum_{k_1=1}^n \sum_{k_2=1}^n \cdots \sum_{k_{\ell}=1}^n P_{k_0, k_1} P_{k_1, k_2} \cdots P_{k_{\ell-1}, k_{\ell}} w_{k_0, k_1} w_{k_1, k_2} \cdots w_{k_{\ell-1}, k_{\ell}} f_{k_{\ell}}. \quad (1.5)$$

1.2 Adjoint Method

A second Monte Carlo method can be derived by considering the linear system adjoint to (1.2)

$$(P^{-1}A)^T \mathbf{y} = \mathbf{d}, \quad (1.6)$$

where \mathbf{y} and \mathbf{d} are the adjoint solution and source term.

By reformulating the fixed point scheme, introducing initial probability, transition probability and weights sequences in the same fashion as done before, the expected value for the estimator becomes

$$\theta_j = E \left[\sum_{\ell=0}^{\infty} W_{\ell} \delta_{k_{\ell}, j} \right] = \sum_{\ell=0}^{\infty} \sum_{k_1=1}^n \sum_{k_2=1}^n \cdots \sum_{k_{\ell}=1}^n b_{k_0} P_{k_0, k_1} P_{k_1, k_2} \cdots P_{k_{\ell-1}, k_{\ell}} w_{k_0, k_1} \cdots w_{k_{\ell-1}, k_{\ell}} \delta_{k_{\ell}, j}. \quad (1.7)$$

We can notice now that the component of the source vector, in each permutation of the random walk, is associated with the initial step. The delta at the end of the series stands for a filter. It means that if I want to estimate the j -th component of the solution vector, just a subset of all the random walks gives a contribution to this. In particular, just those random walks that currently reside on the index I want to estimate are considered.

Chapter 2

Statistical requirements

For the sake of completeness and intelligibility of what will be introduced later on, we report here the statement of the Central Limit Theorem.

Theorem 1 Central Limit Theorem (CLT) for independent sequences

Suppose X_1, X_2, \dots is a sequence of independent identically distributed (i.i.d.) random variables with $E[X_i] = \mu$ and $\text{Var}[X_i] = \sigma^2 < \infty$. Then as $n \rightarrow \infty$, the random variables $\sqrt{n}(S_n - \mu)$ converge in distribution to a normal $\mathcal{N}(0, \sigma^2)$:

$$\sqrt{n} \left(\left(\frac{1}{n} \sum_{i=1}^n X_i \right) - \mu \right) \xrightarrow{d} \mathcal{N}(0, \sigma^2).$$

This is the reference point for what is introduced from now on. In particular the goal is to define a proper transition probability, such that the hypotheses of the theorem hold.

2.1 Forward method - Expected value and variance of the estimator

As already described in Section 1.1 of Report 1, the Forward method is characterized by the introduction of proper weight W_i defined recursively in terms of the steps of the random walk. If these weights get too small, there is no point in keeping the random walk moving, so it is possible to truncate it without significant loss of accuracy.

The formula for the estimator of the Forward Monte Carlo for a generic entry of the solution vector has already been defined in (1.5).

As regards the variance, we know that the following relation holds:

$$\text{Var} \left[\sum_{\ell=0}^{\infty} W_{\ell} f_{k_{\ell}} \right] = E \left[\sum_{\ell=0}^{\infty} W_{\ell}^2 f_{k_{\ell}}^2 \right] - \left(E \left[\sum_{\ell=0}^{\infty} W_{\ell} f_{k_{\ell}} \right] \right)^2. \quad (2.1)$$

Hence the variance can be computed as the difference between the second moment of the random variable and the square of the first moment.

In order to apply the CLT to the estimators defined above, we must require that the estimators have finite expected value and finite variance. This is equivalent in checking the finiteness of the expected value and of the second moment. Therefore we have to impose the following conditions:

$$E \left[\sum_{\ell=0}^{\infty} W_{\ell} f_{k_{\ell}} \right] < \infty \quad (2.2)$$

$$E \left[\sum_{\ell=0}^{\infty} W_{\ell}^2 f_{k_{\ell}}^2 \right] < \infty \quad (2.3)$$

2.2 Adjoint method - Expected value and variance of the estimator

As concerns the Adjoint method the formula for the expected value of the estimator is represented in (1.7).

The variance is

$$\text{Var} \left[\sum_{\ell=0}^{\infty} W_{\ell} f_{k_0} \delta_{k_{\ell},j} \right] = E \left[\sum_{\ell=0}^{\infty} W_{\ell}^2 f_{k_0}^2 \delta_{k_{\ell},j} \right] - \left(E \left[\sum_{\ell=0}^{\infty} W_{\ell} f_{k_0} \delta_{k_{\ell},j} \right] \right)^2 \quad j = 1, \dots, n. \quad (2.4)$$

On the same lead of what we did for the Forward method, we have to impose finiteness of the expected value and second moment as well. Therefore we require the following

$$E \left[\sum_{\ell=0}^{\infty} W_{\ell} \delta_{k_{\ell},j} \right] < \infty \quad j = 1, \dots, n \quad (2.5)$$

and

$$E \left[\sum_{\ell=0}^{\infty} W_{\ell}^2 f_{k_0}^2 \delta_{k_{\ell},j} \right] < \infty \quad j = 1, \dots, n. \quad (2.6)$$

2.3 Numerical interpretation of the statistical requirements

The requests imposed on the expected value and second moment of the estimators can be reformulated in a purely deterministic setting by looking at the spectral radius of some properly defined matrices.

For instance the condition of finiteness of the expected value can be reformulated by requiring

$$\rho(H) < 1, \quad (2.7)$$

where H is the iteration matrix of the fixed point scheme. In fact both equations 1.5 and 1.7 can be considered as power series in terms of H . Moreover having a spectral radius of H smaller than 1 is a necessary and sufficient condition for the Neumann series to converge.

By analogy we aim at reproducing a similar reasoning for the second moment of the estimator, since we want it to be finite as well. By looking at the equations of the variance 2.1 and 2.4 for the Forward and the Adjoint method respectively, we notice that the second moment can be reinterpreted as a power series with respect to the matrices defined as follows:

$$H_{i,j}^* = \frac{H_{i,j}^2}{P_{i,j}} \quad \text{- Forward Method.}$$

or

$$H_{i,j}^* = \frac{H_{j,i}^2}{P_{i,j}} \quad \text{- Adjoint Method.}$$

Accordingly to the necessary and sufficient condition for the power series to converge, we require

$$\rho(H^*) < 1. \quad (2.8)$$

Therefore condition 2.7 and 2.8 are the numerical reformulation of the finiteness of the expected value and second moment of the random variables considered. These results can be found in [HMY13].

2.4 Necessary conditions, sufficient conditions and choice of the transition matrix

In this chapter we get down to enunciate some results presented in [HMY13] about necessary conditions and sufficient conditions for convergence. In particular we discuss what is a suitable choice for constructing a transition probability that makes the MC solver converge.

We remember that the construction of the transition probability must satisfy some constraints such as

$$\begin{cases} P_{i,j} \geq 0 \\ \sum_{j=1}^N P_{i,j} = 1 \end{cases}$$

plus one more restriction which varies depending on the particular Monte Carlo approach chosen

Forward Method: $H_{i,j} \neq 0 \Rightarrow P_{i,j} \neq 0$

Adjoint Method: $H_{j,i} \neq 0 \Rightarrow P_{i,j} \neq 0$

The previous restrictions on the attainable values for the entries of the transition probability are called *transition conditions*.

We start by presenting

Theorem 2 Consider a vector $a = (a_1, a_2, \dots, a_N)^T$ where at least one element is nonzero, $a_k \neq 0$ for some $k \in \{1, \dots, N\}$.

- For a probability vector $p = (p_1, p_2, \dots, p_N)^T$ satisfying the transition conditions, the lower bound of $\sum_{k=0}^N \frac{a_k^2}{p_k}$ is $\left(\sum_{k=1}^N |a_k| \right)^2$
- There always exists a probability vector such that $\sum_{k=0}^N \frac{a_k^2}{p_k} \geq c \geq 1$ for all $c > 1$.

It might be observed that the minimum of the quantity $\sum_{k=0}^N \frac{a_k^2}{p_k}$ is attained when the probability vector is defined such as $p_k = \frac{|a_k|}{\sum_{k=1}^N |a_k|}$. It implies that the MAO (Monte Carlo Almost Optimal) probability is the one that reduces as much as possible the ∞ -norm of H^* .

For the sake of simplicity let us introduce now a generic random walk truncated at a certain k -th step

$$\gamma_k : r_0 \rightarrow r_1 \rightarrow r_2 \rightarrow \dots \rightarrow r_k$$

and the statistical estimator

$$X(\gamma_k) = \frac{H_{r_0, r_1} H_{r_1, r_2} \dots H_{r_{k-1}, r_k}}{P_{r_0, r_1} P_{r_1, r_2} \dots P_{r_{k-1}, r_k}} f_{r_k}$$

which is associated with the Forward Method.

Then the following theorem holds

Theorem 3 (Suited to the Forward Method)

Let $H \in \mathbb{R}^{n \times n}$ such that $\|H\|_\infty < 1$. Consider ν_k as the realization of a random walk γ truncated at the k -th step. Then, there always exists a transition matrix P such that $\text{Var}(X(\nu_k)) \rightarrow 0$ and $\text{Var}\left(\sum_{\nu} X(\nu_k)\right)$ is bounded as $k \rightarrow \infty$.

If we introduce the estimator

$$X(\gamma_k) = \frac{H_{r_0, r_1}^T H_{r_1, r_2}^T \cdots H_{r_{k-1}, r_k}^T \text{sign}(f_{r_0}) \|\mathbf{f}\|_1}{P_{r_0, r_1} P_{r_1, r_2} \cdots P_{r_{k-1}, r_k}}$$

which is associated with the Adjoint Method, then we can state a theorem specular to 3:

Theorem 4 (*Suited to the Adjoint Method*)

Let $H \in \mathbb{R}^{n \times n}$ with $\|H\|_1 < 1$. Consider ν_k as the realization of a random walk γ truncated at the k -th step. Then, there always exists a transition matrix P such that $\text{Var}(X(\nu_k)) \rightarrow 0$ and $\text{Var}(\sum_{\nu} X(\nu_k))$ is bounded as $k \rightarrow \infty$.

In particular, relying on what guaranteed by Theorem 2, in both Theorems 3 and 4 the condition holds for sure with the MAO transition matrix.

These theoretical results represent sufficient condition for the Forward and Adjoint Monte Carlo to converge and can be easily checked. However requiring $\|H\|_{\infty} < 1$ rather than $\|H\|_1 < 1$ may be too demanding. There exist some preconditioners that guarantee that this condition be respected for sure (e.g. approximate inverse preconditioners with a proper value of dropping tolerance). Nevertheless it may require the loss of sparsity of the iteration matrix and of course this is something we cannot accept in our numerical setting.

Therefore, in situation where 3 and 4 are too strong conditions to be verified, the necessary and sufficient condition on the spectral radius of H^* is the only criterion to look at. The drawback of the necessary and sufficient condition, however, consists in computing the spectral radius of a matrix. Indeed the computational cost equals the one of a deterministic solver for the resolution of the original linear system we are interested in.

We now provide a condition that, if verified, entails the impossibility of reaching convergence with any Monte Carlo linear solver.

Theorem 5 Let H be an $N \times N$ matrix with spectral radius $\rho(H) < 1$. Let H^+ be the $N \times N$ matrix where $H_{i,j}^+ = |H_{i,j}|$. If $\rho(H^+) > 1$, there does not exist a transition matrix P satisfying the transition conditions such that the variance $\text{Var}(X(\gamma_k))$ converges to zero as $k \rightarrow \infty$.

Theorem 5 hold for both Forward and Adjoint method, since H^+ and $(H^+)^T$ have the same spectral radius.

Chapter 3

Adaptive approaches

3.1 An adaptive cut-off for the length of random walks

At first we focus on the definition of a weight cut-off in order to decide where to terminate the histories. It implies that we are looking for a quantity that tells us how many terms of the Neumann series need to be considered.

The goal is to find an automatic way for the cutoff of the histories. Therefore we are looking for an m such that

$$\tilde{\theta}_i = E \left[\sum_{\ell=0}^m W_{\ell} f_{k_{\ell}} \right] = x_i = \sum_{\ell=0}^m \sum_{k_1=1}^n \sum_{k_2=1}^n \cdots \sum_{k_{\ell}=1}^n P_{k_0, k_1} P_{k_1, k_2} \cdots P_{k_{\ell-1}, k_{\ell}} w_{k_0, k_1} w_{k_1, k_2} \cdots w_{k_{\ell-1}, k_{\ell}} f_{k_{\ell}}.$$

and

$$\tilde{\theta}_j = E \left[\sum_{\ell=0}^m W_{\ell} \delta_{k_{\ell}, j} \right] = \sum_{\ell=0}^m \sum_{k_1=1}^n \sum_{k_2=1}^n \cdots \sum_{k_{\ell}=1}^n f_{k_0} P_{k_0, k_1} P_{k_1, k_2} \cdots P_{k_{\ell-1}, k_{\ell}} w_{k_0, k_1} \cdots w_{k_{\ell-1}, k_{\ell}} \delta_{k_{\ell}, j}.$$

are good approximations of 1.5 and 1.7 respectively.

In [Sla13] there is a criterion for the cutoff of the random walk that works either for the Forward or for the Adjoint method. It consists in defining a cutoff relative threshold W_c and looking for a step such that

$$W_m \leq W_c W_0. \quad (3.1)$$

W_0 is the value of the weight at the initial step of the random walk and W_m is the value of the weight after m steps.

3.1.1 Relative standard deviation - Forward method

As regards the Forward method, we know that the expression for the variance is defined in Formula (2.1).

In this context, a reasonable criterion to determine the number \tilde{N}_i of random walks to be run is setting a threshold ε_1 and determine

$$\tilde{N}_i \quad s.t. \quad \frac{\sqrt{Var[\theta_i]}}{|E[\theta_i]|} < \varepsilon_1, \quad i = 1, \dots, n. \quad (3.2)$$

The dependence of $Var[\theta_i]$ and $E[\theta_i]$ on \tilde{N}_i , which seems to be absent in the previous formula, is highlighted by the fact that θ_i is estimated by fixing a finite number of histories. Therefore we are controlling the relative standard deviation requiring it not to be too large. In other words we are pursuing a statistical setting where the uncertainty factor is not dominating over the expected value. This simple adaptive approach can be applied for the estimation of each component x_i . Therefore different number of histories may be employed to compute different entries of the solution vector.

3.1.2 Relative standard deviation - Adjoint method

As we already stressed out in the previous reports, in the Adjoint method each random walk gives contributions for more than one entry. Therefore in this case the selection of the number of random walks is global, since it involves the total number of histories to estimate the entire solution vector.

As concerns the Adjoint method, the estimation of the variance for each entry is represented by Formula (2.4).

A possible adaptive selection of \tilde{N} , in this situation, is

$$\tilde{N} \quad s.t. \quad \frac{\|\sigma_{\tilde{N}}\|_1}{\|\mathbf{x}\|_1} < \varepsilon_1, \quad (3.3)$$

where σ is a vector whose entries are $\sigma_{\tilde{N},i} = Var[\theta_i]$.

3.1.3 An a posteriori variance-based adaptivity

What introduced in the two previous section can be exploited in order to build an a posteriori adaptive algorithm, capable to identify the minimal value of \tilde{N} that verifies 3.2 or 3.3 respectively.

Below there are the pseudo-codes associated with both Forward and Adjoint Monte Carlo.

Data: N, ε_1

Result: $\tilde{N}_i, \sigma_i, x_i$

For each entry of the solution vector ;

$\tilde{N}_i = N$;

compute θ_i ;

while $\frac{\sigma_i}{|E[\theta_i]|} < \varepsilon_1$ **do**

$\tilde{N} = \tilde{N} + N$;

 compute $E[\theta_i] = x_i$;

end

return θ_i, x_i, σ_i ;

Algorithm 1: A posteriori adaptive Forward Monte Carlo

Data: N, ε_1

Result: $\tilde{N}, \sigma_i, \mathbf{x}$

$\tilde{N} = N$;

compute θ ;

while $\frac{\|\sigma\|}{\|E[\theta]\|} < \varepsilon_1$ **do**

$\tilde{N} = \tilde{N} + N$;

 compute $E[\theta] = \mathbf{x}$;

end

return $\theta, \mathbf{x}, \sigma$;

Algorithm 2: A posteriori adaptive Adjoint Monte Carlo

3.2 Numerical results

The adaptive threshold is set such that $\varepsilon_1 = 0.01$.

Forward Monte Carlo

As concerns the Forward method, the maximal admitted number of histories for each component is equal to 10 times the length of the solution vector. The number of steps per random walk is initially fixed to a constant number equal to 1000. At each adaptive iteration the number of random walks employed is increased by ten. Even if this is a very little number and it compromises the efficiency of the algorithms in terms of speed, we want to keep it small. In fact for now we want to figure out what is the minimal number of histories required for the fulfillment of the adaptive criterion. In Table 3.1 there are results just for the one-dimensional shifted Laplacian, the "JPWH_991" and the thermal equation. Result for the "FS_680_1" are missing, since $\rho(H^*) = 1.2554$ implies impossibility to gain convergence. Instead as concerns the "ifiss_advection_diffusion" and the two-dimensional Laplacian results are not shown because the adaptive criterion blows up the number of random walks required.

Type of problem	Relative Error	CPU Time (s)	Nb. Histories
1d shifted Laplacian	0.0174	0.43	1324
2d Laplacian	-	-	-
JPWH_991	0.011	18058	6739640
Marshak Equation	0.0175	543	5238
ifiss	-	-	-

Table 3.1: Forward Monte Carlo. Adaptive criterion.

The algorithms to detect the number of histories to run and the history's length cutoff can be combined in a unique adaptive approach. The threshold for the weight cut-off is set to 10^{-6} . Results for the "ifiss_advection_diffusion" are missing for the same reason as above.

Type of problem	Relative Error	CPU Time (s)	Nb. Histories
1d shifted Laplacian	0.0145	5.26	1626
2d Laplacian	0.333	210.8	9000
ifiss advection diffusion			
JPWH_991	0.29	24.37	9910
Marshak Equation	0.017	3.54	5172

Table 3.2: Forward Monte Carlo. Adaptive cut-off and adaptive selection of # histories.

Adjoint Monte Carlo

Now we get down to the Adjoint method. At each adaptive iteration the number of random walks employed is increased by two.

No restraints are put for a maximal number of random walks admitted. therefore we let the algorithm free to run until the constraint is verified.

Here below Table 3.3 shows results for the different test cases considered. As we can see, the times employed are very high. However we have to note that this is the standard Monte Carlo algorithm whose rate of convergence is the square root of the random walks computed. Thus we know it to be very slow a priori. Moreover the request of having a relative standard deviation smaller than 0.1 corresponds to a relative error of the same order. Therefore results are coherent with what expected a priori. For "JPWH_991" and "FS_680_1" the Adjoint method is not a convergent method, since $\rho(H^*) = 1.0505$ and $\rho(H^*) = 3.598$ respectively.

Type of problem	Relative Error	CPU Time (s)	Nb. Histories
1d shifted Laplacian	0.139	4.19	730
2d Laplacian	0.136	233.4	570
ifiss advection diffusion	0.282	123.7	630
Marshak equation	0.288	369	880

Table 3.3: Adjoint Monte Carlo. Adaptive criterion.

3.3 A stationarity check for the a posteriori adaptivity

The a posteriori approach introduced before is very naive and it does not take into account whether the number of histories is enough to reach the asymptotic trend or not. In order to do this we need to look at the value of the variance with respect to two consecutive adaptive steps. For the Forward method it means that we need to set the number of histories \tilde{N}_i such that

$$\frac{|\sigma_i^{\tilde{N}_i-1} - \sigma_i^{\tilde{N}_i}|}{|x_i|} < \varepsilon_2 \quad i = 1, \dots, n \quad (3.4)$$

where n is the size of the problem and ε_2 is a threshold set for the check of stationarity.

As concerns the Adjoint method instead we need to pursue \tilde{N} such that

$$\frac{\|\boldsymbol{\sigma}^{\tilde{N}-1} - \boldsymbol{\sigma}^{\tilde{N}}\|_1}{\|\mathbf{x}\|_1} < \varepsilon_2. \quad (3.5)$$

3.4 can be combined with 3.2 and 3.5 can be combined with 3.3. In this way we can build up a more robust adaptive approach that takes into account the relative magnitude of the variance and the achievement of the asymptotic behavior as well. However it is necessary to find a compromise between the accuracy expected and the computational cost. Therefore in practical situations the stationarity check may require too effort to be employed.

3.4 An a priori adaptive approach on the residual

An appealing alternative might be the use of an a priori error estimator. In order to formulate such this device, assume θ is an estimator for a quantity x which is unknown. It holds

$$Var[\theta] = E[\theta^2] - (E[\theta])^2.$$

The theoretical variance can be estimated with a sample quantity such as the sample variance, defined as

$$\hat{\sigma} = \frac{1}{(n-1)} \sum_{i=1}^n (\theta_i - \bar{\theta})^2, \quad \bar{\theta} = \frac{1}{n} \sum_{i=1}^n \theta_i. \quad (3.6)$$

Each θ_i is a random variable associated with a single random walk. The expected value of all the random walks return the estimator we are interested in. Therefore the sample size n is the total number of histories run. Since all the random walks are independent one of the other. Therefore all the realizations of the estimator θ_i are independent as well.

This means that

$$cov(\theta_i, \theta_j) = 0, \quad i \neq j.$$

This consideration is very important for our applications. In fact it implies that the standard deviation and the statistical error both associated with the standard Monte Carlo scales with the same order.

The statistical error is defined as

$$err = \sqrt{\frac{1}{n} \left[\sum_{i=1}^n (\theta_i - x)^2 \right]}. \quad (3.7)$$

However we are assuming that θ be a unbiased estimator. This entails that statistical error and standard deviation are actually the same thing.

By combining 3.6 and 3.7 we can conclude that

$$\frac{\|\mathbf{x} - \hat{\mathbf{x}}\|}{\|\mathbf{x}\|} \sim \frac{C}{\sqrt{n}}$$

for some constant $C > 0$ and $\hat{\mathbf{x}}$ represents the standard Monte Carlo solution.

Given a linear system

$$A\mathbf{x} = \mathbf{b}, \quad A \in \mathbb{R}^{d \times d}, \quad \mathbf{x}, \mathbf{b} \in \mathbb{R}^d$$

and given $\hat{\mathbf{x}}$ the standard Monte Carlo solution, the residual is defined such as

$$\mathbf{r} = \mathbf{b} - A\hat{\mathbf{x}} = A(\mathbf{x} - \hat{\mathbf{x}}).$$

It entails that the residual has the same behavior as the error. In fact the former is a linear transformation of the latter.

Hence the following relationship holds

$$\log(\|\mathbf{r}\|_2) = -\frac{1}{2} \log(n) + \log(\|\mathbf{b}\|_2) \quad (3.8)$$

regardless of the basis chosen for the logarithm. 3.8 is equivalent to

$$\log(\|\mathbf{r}\|_2) = \log\left(\frac{\|\mathbf{b}\|_2}{\sqrt{n}}\right).$$

Thus the following relationship between the residual after n histories and the right hand side holds:

$$\|\mathbf{r}\|_2 = \frac{\|\mathbf{b}\|_2}{\sqrt{n}}.$$

Since it is known that

$$\frac{\|\mathbf{x} - \hat{\mathbf{x}}\|_2}{\|\mathbf{x}\|_2} \leq K_2(A) \frac{\|\mathbf{r}\|_2}{\|\mathbf{b}\|_2} \quad (3.9)$$

By imposing

$$K_2(A) \frac{\|\mathbf{r}\|}{\|\mathbf{b}\|} < \varepsilon, \quad 0 < \varepsilon < 1$$

we get

$$K_2(A) \frac{\|\mathbf{b}\|}{\sqrt{n}} < \varepsilon \|\mathbf{b}\|, \quad 0 < \varepsilon < 1$$

and it leads to

$$n > \frac{\left(K_2(A)\right)^2}{\varepsilon^2}. \quad (3.10)$$

3.10 represents an estimation of the number of random walks to run in order to have the guarantee that the relative error is below a certain threshold ε . Of course 3.9 may not be an effective estimation of the relative error in case of large values of the condition number. In this case the quality of 3.10 is affected as well.

Chapter 4

Linear solvers with hybrid schemes

Even if the standard Monte Carlo, in both its Forward and Adjoint formulation, has the advantage of being embarrassingly parallelizable, it has a drawback as well. This is its slow rate of convergence, associated with the $\frac{1}{\sqrt{N}}$ behavior predicted by the CLT.

Moreover, when the spectral radius of the iteration matrix is near the unit value, Monte Carlo direct methods reach convergence after many steps for each random walk. Therefore the standard Monte Carlo paradigm is most of the time unpractical for solving linear systems associated with real applications.

4.1 Monte Carlo Synthetic Acceleration

Recently Evans et al. introduced a new method: the *Monte Carlo Synthetic Acceleration (MCSA)* ([ESW13] and [EMSH14]). If we think about a fixed point formulation of the problem such as $\mathbf{x} = H\mathbf{x} + \mathbf{b}$, then the Monte Carlo Synthetic Acceleration assumes this form

```

Data:  $H, \mathbf{b}, \mathbf{x}_0$ 
Result:  $x_{num}$ 
 $\mathbf{x}^l = \mathbf{x}_0$ ;
while not reached convergence do
     $\mathbf{x}^{l+\frac{1}{2}} = H\mathbf{x}^l + \mathbf{b}$ ;
     $\mathbf{r}^{l+\frac{1}{2}} = \mathbf{b} - A\mathbf{x}^{l+\frac{1}{2}}$ ;
     $\delta\mathbf{x}^{l+\frac{1}{2}} = (I - H)^{-1}\mathbf{r}^{l+\frac{1}{2}}$ ;
     $\mathbf{x}^{l+1} = \mathbf{x}^{l+\frac{1}{2}} + \delta\mathbf{x}^{l+\frac{1}{2}}$ ;
end
 $x_{num} = x^{l+1}$ ;

```

Algorithm 3: Monte Carlo Synthetic Acceleration

The Monte Carlo method is used to compute the updating contribution $\delta\mathbf{x}^{l+\frac{1}{2}}$.

4.2 Preconditioning techniques

In this section we focus on the pursuit of preconditioners that facilitate the convergence of the stochastic schemes introduced so far.

4.2.1 Block Preconditioning

Consider a matrix $A \in \mathbb{R}^{n \times n}$ of the form

$$A = \begin{bmatrix} A_{11} & A_{12} & \cdots & A_{1p} \\ A_{21} & A_{22} & \cdots & A_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ A_{p1} & \cdots & \cdots & A_{pp} \end{bmatrix}$$

where p is a divisor of n and such that $A_{ii} \in \mathbb{R}^{n_i \times n_i}$ are nonsingular square matrices. By defining a block diagonal matrix such as

$$D = \begin{bmatrix} A_{11} & 0_{12} & \cdots & 0_{1p} \\ 0_{21} & A_{22} & \cdots & 0_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ 0_{p1} & \cdots & \cdots & A_{pp} \end{bmatrix}$$

it is possible to compute its inverse. D^{-1} can be used as a preconditioner for the original matrix A . In particular

$$D^{-1}A = \begin{bmatrix} I_{n_1 \times n_1} & A_{11}^{-1}A_{12} & \cdots & \cdots & A_{11}^{-1}A_{1p} \\ A_{22}^{-1}A_{21} & I_{n_2 \times n_2} & A_{22}^{-1}A_{23} & \cdots & A_{22}^{-1}A_{2p} \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ A_{pp}^{-1}A_{p1} & \cdots & \cdots & A_{pp}^{-1}A_{p(p-1)} & I_{n_p \times n_p} \end{bmatrix}.$$

The ultimate goal is to resort to this preconditioning technique in order to achieve the condition $\rho(D^{-1}A) < 1$. In fact this is necessary in order to reformulate 1.1 into 1.3.

Convergence of 1.3 with block diagonal preconditioning, from the deterministic point of view, is guaranteed when A is an irreducibly diagonally dominant matrix or when A is an M -matrix. In fact in situation such as these, theoretical results guarantee that $A = D - N$ is a regular splitting (see [Axe96]).

4.2.2 Approximate Inverse of a Sparse Matrix (AINV)

This stimulated our interests into other kind of preconditioners. At first we concentrated on Approximate Inverse preconditioners (see [Saa03] and [Ben02]).

The idea is to search an approximate inverse of an arbitrary sparse matrix by minimizing the Frobenius norm of the residual matrix $I - AM$:

$$F(M) = \|I - AM\|_F^2. \quad (4.1)$$

A matrix M whose value $F(M)$ is small is defined as a right-looking approximate inverse of A . In a similar way it is possible to define a left-looking approximate inverse such the the following objective function is minimized:

$$\|I - MA\|_F^2. \quad (4.2)$$

A further alternative may imply to pursuit a left-right pair L, U such that

$$\|I - LAU\|_F^2. \quad (4.3)$$

4.1 is equivalent to solve a set of decoupled least square problems:

$$F(M) = \|I - AM\|_F^2 = \sum_{j=1}^n \|e_j - Am_j\|_2^2, \quad (4.4)$$

where e_j and m_j are the j -th columns of the identity matrix and of the matrix M . An approach that may be used to solve 4.4 is to minimize separately each term of the sum in such this way:

$$f_j(m) = \|e_j - Am_j\|_2^2, \quad j = 1, 2, \dots, n.$$

There are many practical techniques that may be applied in order to build the AINV preconditioner. One of these is the one described in [BMM96] and [BT98]. This is the actual approach we applied to the construction of the preconditioner. The specific name associated with this approach is *Factored Inverses via Orthogonalization*. The idea is to compute an approximate factorization of the form $W^T AZ = D$. W and Z are unit upper triangular matrices and D is a diagonal matrix.

In order to compute the AINV preconditioner a FORTRAN code provided by Miroslav Tuma (Academy of Sciences, Prague) is used. We report as follows the algorithm employed for the construction of the so called *Right-looking Factored AINV*. The algorithm provide the user with a

dropping tolerance which can be used in order to tune the fill-in effect of the factored preconditioner.

Data: $p = q = (0, \dots, 0) \in \mathbb{R}^n$, $Z = [z_1, \dots, z_n] = I_n$, $W = [w_1, \dots, w_n] = I_n$

Result: Z, W

```

for  $k = 1, \dots, n$  do
     $p_k = w_k^T A e_k$ ,  $q_k = e_k^T A z_k$ ;
    for  $i = k + 1, \dots, n$  do
         $p_i = (w_k^T A e_k) / p_k$ ,  $q_i = (e_k^T A z_i) / q_k$ ;
        Apply a dropping rule to  $p_i, q_i$ ;
         $w_i = w_i - w_k p_i$ ,  $z_i = z_i - z_k q_i$ 
        Apply a dropping rule to  $w_{j,i}$  and  $z_{j,i}$  for  $j = 1, \dots, i$ ;
    end
    ;
end
;
Choose diagonal entries of  $D$  as the component of  $p$  or  $q$ ;
return  $Z, W$ ;

```

Algorithm 4: Right-looking Factored AINV

4.2.3 Incomplete LU Factorization Preconditioners (ILU)

This set of preconditioners has the goal of computing a sparse lower triangular matrix L and a sparse upper triangular matrix U such that the residual matrix $R = LU - A$ respect some constraints. One of these constraints may be the fact that some entries must be set to zero. For instance we can introduce a zero pattern set

$$P \subset \{(i, j) | i \neq j; 1 \leq i, j \leq n\}$$

and use this set in order to impose some constraints on the shape of the preconditioner built.

Different variants of the Incomplete LU Factorization algorithm are available (see [?]). However it is possible to prove that if all the formulation are well defined, then they are all equivalent. One of these variants is the following

```

Data:  $P$ 
Result:  $A$ 
For each  $(i, j) \in P$  set  $a_{ij} = 0$ ;
for  $k = 1, \dots, n - 1$  do
    for  $i = k + 1, n$  do
        if  $(i, k) \notin P$  then
             $a_{ik} = \frac{a_{ik}}{a_{kk}}$  for  $j = k + 1, \dots, n$  and for  $(i, j) \notin P$  do
                 $a_{ij} = a_{ij} - a_{ik} * a_{kj}$ 
            end
        end
    end
end
return  $Z, W$ ;

```

Algorithm 5: General Static Pattern ILU

This kind of algorithm is guaranteed to terminate without any breakdown just for M -matrices. Assuming that the zero pattern P coincides with the zero pattern of A leads to the $ILU(0)$.

```

Data:  $P$ 
Result:  $A$ 
for  $i = 2, \dots, n$  do
  for  $k = i, i - 1$  and for  $(i, k) \in NZ(A)$  do
     $a_{ik} = \frac{a_{ik}}{a_{kk}}$  for  $j = k + 1, \dots, n$  and for  $(i, j) \in NZ(A)$  do
       $a_{ij} = a_{ij} - a_{ik}a_{kj}$ 
    end
  end
end

```

Algorithm 6: ILU(0).

More sophisticated variants of the algorithm enable to enlighten the constraint of the sparsity pattern, finding a compromise between sparsity and accuracy in approximating the inverse of A . We refer to [Saa03] and [Ben02] for further information about it.

4.3 Reordering techniques

Reordering techniques are often used in order to improve the performance of preconditioners computed through a factorization approach. The reason why they are employed is to reduce the fill-in effect, especially for ILU algorithms.

For instance, if P and Q are permutation matrices, system 1.1 may be turned into the following

$$PAQ\mathbf{y} = P\mathbf{b}, \quad \mathbf{x} = Q\mathbf{y}.$$

In general it is not possible to find a relationship between the shape of factored preconditioners associated with PAQ with respect to the preconditioners of A .

A common reordering employed for a generic matrix is the *Reverse Cuthill McKee ordering* (RCM). The purpose of this reordering is to reduce the band of nonzero entries. This way factored preconditioners such as the ones computed via ILU are expected to preserve the same bandwidth. Effectiveness of this way of proceeding in many applications is proved in [Ben02].

4.4 Numerical results

4.4.1 Illustrative examples

In this section we apply the Sequential Monte Carlo and the Monte Carlo synthetic Acceleration on the same test cases we have considered in the previous chapters. We use an adaptive approach to truncate the random walk at each numerical iteration by using a weight cut-off $W_c = 10^{-6}$. For both the Forward and the Adjoint method we use an adaptive approach also to decide how many random walks to employ at each numerical iteration. In particular as concerns the Forward method we add 10 histories at each adaptive check, whereas for the Adjoint method we compute 100 histories at each adaptive check.

The preconditioner applied on the linear systems is still a diagonal one, as for all the tests run in the previous chapters as well.

Type of problem	Rel. Residual	Rel. Error	Numerical Iter.	CPU Time (s)
1d shifted Laplacian	$1.99 \cdot 10^{-8}$	$8.64 \cdot 10^{-9}$	6	1.88
2d Laplacian	$9.54 \cdot 10^{-8}$	$4.027 \cdot 10^{-8}$	275	910
ifiss advection diffusion	$9.24 \cdot 10^{-8}$	$1.29 \cdot 10^{-6}$	73	808
JPWH 991	$8.41 \cdot 10^{-8}$	$5.21 \cdot 10^{-8}$	67	409
Marshak Equation	$1.51 \cdot 10^{-8}$	$5.68 \cdot 10^{-4}$	7	73.1

Table 4.1: Forward MCSA.

4.4.2 SP_N equations

In this subsection we focus on a set of partial differential equations which are considered a milestone in nuclear physics: the SP_N equations. These kind of equations is employed, for instance, to

Type of problem	Rel. Residual	Rel. Error	Numerical Iter.	CPU Time (s)
1d shifted Laplacian	$1.99 \cdot 10^{-8}$	$8.64 \cdot 10^{-9}$	6	1.88
2d Laplacian	$9.54 \cdot 10^{-8}$	$4.027 \cdot 10^{-8}$	275	910
ifiss advection diffusion	$9.24 \cdot 10^{-8}$	$1.29 \cdot 10^{-6}$	73	808
Marshak Equation	$1.51 \cdot 10^{-8}$	$5.68 \cdot 10^{-4}$	7	73.1

Table 4.2: Adjoint MCSA.

synthesize phenomena such as the power distribution cross a nuclear reactor core (see [EMSH14]). Physical events of this type are extremely complicated and practical devices able to analyze them in their entire complexity are not at our hand currently. Indeed the starting equation, used as a reference to represent the evolution of phenomena such as these, is the Boltzmann equation. In a multidimensional domain the terms considered inside of the equation are difficult to be treated and state-of-the art modelings are not capable of handling the complexity of the original model. Therefore successive simplifications are introduced in order to formulate a treatable and practical representation of the starting problem.

For example, let us consider a steady-state, multi-group, one-dimensional, eigenvalue-form of Boltzmann:

$$\begin{aligned} \mu \frac{\partial \psi^g(x, \mu)}{\partial x} + \sigma^g(x) \psi^g(x, \mu) = & \sum_{g'=1}^{N_g} \int_{4\pi} \sigma_s^{gg'}(x, \hat{\Omega} \cdot \hat{\Omega}') \psi^{g'}(x, \Omega') d\Omega' + \\ & + \frac{1}{k} \sum_{g'=1}^{N_g} \frac{\chi^g}{4\pi} \int_{4\pi} \nu \sigma_f^{g'}(x) \psi^{g'}(x, \Omega') d\Omega'. \end{aligned} \quad (4.5)$$

The meaning of the quantities appearing in 4.5 is the following:

- N_g = total number of levels of energy used to classify neutrons
- $\psi^g(x, \mu)$ = angular flux for group g
- σ^g = total interaction cross section
- $\sigma_s^{gg'}(x, \hat{\Omega} \cdot \hat{\Omega}')$ = scattering cross section from group $g' \rightarrow g$
- χ^g = resulting fission spectrum in group g
- k = ratio of neutron populations in subsequent fission generations.

By resorting to a spectral decomposition of the angular flux ψ and of the interaction σ by Legendre polynomials we get

$$\psi^g(\mu) = \sum_{i=0}^{\infty} \frac{2i+1}{4\pi} \varphi_i^g p_i(x)$$

and


$$\sigma_s(\mu_0) = \sum_{j=0}^{\infty} \frac{2j+1}{4\pi} \sigma_{sj} p_j(x).$$

The truncation of the Fourier series to the N -th term we get to the formulation of the so called P_N equations:

$$\frac{\partial}{\partial x} \left[\frac{i}{2i+1} \varphi_{i-1}^g + \frac{i+1}{2i+1} \varphi_{i+1}^g \right] + \sum_{g'=1}^{N_g} (\sigma^g \delta_{gg'} - \sigma_{sn}^{gg'}) \varphi_i^{g'} = \frac{1}{k} \sum_{g'=1}^{N_g} \chi^g \nu \sigma_f^{g'} \varphi_i^{g'} \delta_{i0}, \quad i = 0, 1, 2, \dots, N.$$

Considering just odd sets of P_N equations and by removing lower gradient terms from each equation we get to the formulation of simplified P_N equations (SP_N):

$$-\nabla \cdot \mathbb{D}_i \nabla \mathbb{U}_i + \sum_{j=0}^{\frac{N+1}{2}} \mathbb{A}_{ij} \mathbb{U}_j = \frac{1}{k} \sum_{j=1}^{\frac{N+1}{2}} \mathbb{F}_{ij} \mathbb{U}_j, \quad i = 1, \dots, \frac{N+1}{2} \quad (4.6)$$

Figure 4.1: Sparsity pattern of the SP_1 equation.

\mathbb{U}_j is a linear combination of moments.

In the problems we focused on, a Finite Volume discretization has been applied in order to build the linear system to solve.

We consider two examples of SP_1 equations, an example of SP_3 and one of SP_5 . In particular for the SP_{11} equations the matrices have respectively size 18207×18207 and 19941×19941 . Their sparsity pattern is shown in Figures 4.1 and 4.2. The SP_3 matrix is a 36414×36414 matrix while the SP_5 one is a 54621×54621 .

The goal is to detect a suitable preconditioner to apply to the linear system in order to make the MC solver converge. As it will be further explained in the following chapter, having a diagonally dominant structure might be helpful but actually this is not our case. In fact we do not succeed in gaining convergence of the Monte Carlo linear solver with a diagonal preconditioning for these problems.


Therefore a block diagonal preconditioner is initially tested, by varying the size of the blocks in a reasonable range of values. In fact it is very important to preserve the sparsity pattern of the ultimate linear system. It turns out that for the two cases of SP_1 equations we have to apply block diagonal preconditioners of sizes respectively 63×63 and 69×69 in order to get the conditions stated in [HMY13] verified.

The effect of the application of the preconditioner in terms of fill-in effect are represented in Table 4.3. For the matrix (b) the preconditioner is quite efficient since it does not change much the number of nonzero entries in the matrix. The same does not hold for matrix (a), since in this case the preconditioner increases the number of nonzero entries by a factor of six.

Kind of matrix	size	nnz	Prec. block size	nnz after prec
SP_1 (a)	18207	486438	63	2644523
SP_1 (b)	19941	998631	69	1774786

Table 4.3: SP_1 matrices. Analysis of sparsity.

Numerical results are shown in Table 4.4. In terms of number of numerical iterations employed the results are satisfactory. As concerns the times, the slowness of the method is an index of the fact that most of the work is executed by the stochastic part of the algorithm. Therefore most of the computation is relegated into the random walks which in this case are all run in a serial mode.

Figure 4.2: Zoom on the sparsity pattern of the SP_1 equation.

Therefore, even if the times themselves are not satisfactory, they are promising in the viewpoint of a future parallelization of the method.

matrix	$\rho(H)$	$\rho(H^*)$	relative err.	# iterations	CPU time (s)
SP_1 (a)	0.9779	0.9758	$9.97 \cdot 10^{-6}$	340	17035 ($\approx 5h$)
SP_1 (b)	0.9798	0.9439	$3.89 \cdot 10^{-5}$	209	11204 ($\approx 3h$)

Table 4.4: SP_1 matrices. Block diagonal preconditioning. Numerical results

The block diagonal preconditioning is not successful in general on SP_N matrices. In fact by adopting the same approach on SP_3 and SP_5 matrices it was not possible to find a configuration such that the MC solvers would converge. Even by varying the size of the blocks, it is easy to have $\rho(H) < 1$ but the same does not hold for H^* . Both left and right preconditioning have been tried out as well.

We have computed the AINV preconditioner on all the SP_N matrices at hand to pursue a configuration suitable for us. In order to accomplish this, we tuned the drop tolerance that controls the fill-in phenomenon of the preconditioner. Once this done, we have to choose the position where to apply the preconditioner, on the left rather than on the right.

In fact in the deterministic environment this would not make any substantial difference, since the resulting iteration matrix would have the same spectral radius regardless of the position where the preconditioning is applied. However the same does not hold in the definition of the Monte Carlo linear solver. In fact the matrix H^* is defined entry-wise in terms of H . Therefore the position where the preconditioner is applied affects the values attained by the entries of H^* , affecting its spectral radius as well.

In conclusion the position of the preconditioning may determine the convergence of the stochastic scheme.

An illustrative example in this case is represented indeed by the AINV preconditioning over the two SP_1 matrices at hand.

For the $SP_1(a)$ we compute a factored AINV preconditioner with a drop tolerance of $\varepsilon = 0.001$ for both the factors. In Table 4.5 we report the value of the spectral radii of H^* for $SP_1(a)$ and $SP_1(b)$ when the preconditioning is applied on the left and on the right respectively.

Initial matrix	$\rho(H)$	$Leftprec - \rho(H^*)$	$Rightprec - \rho(H^*)$
SP_1 (a)	0.722	1.240	0.920
SP_1 (b)	0.644	1.185	0.998

Table 4.5: Spectral radii of H and H^* for different AINV preconditioning positions.

Computing the entity of the fill-in for the right preconditioning we get

$$\text{ratio} = \frac{nnz(H)}{nnz(A)} = 22.07,$$

where $H = I - AP^{-1}$ and A is the matrix associated with the $SP_1(a)$ problem.

Therefore the fill-in effect compromises the sparsity pattern.

As regards the $SP_2(b)$ matrix an AINV preconditioner with a drop tolerance of $\varepsilon = 0.0009$ for both the factors has been computed. Checking the magnitude of the fill-in phenomenon for the right preconditioning we get

$$\text{ratio} = \frac{nnz(H)}{nnz(A)} = 14.94.$$

In this case we have a massive fill-in as well.

In Table 4.6 we report the numerical results associated with the two aforementioned cases.

matrix	relative err.	# iterations	CPU time (s)
SP_1 (a)	$6.277 \cdot 10^{-7}$	33	900 ($\approx 15\text{min}$)
SP_1 (b)	$9.885 \cdot 10^{-7}$	21	650 ($\approx 10\text{min}$)

Table 4.6: Numerical results for different AINV preconditioning positions.

Both in terms of numerical iterations and time employed the AINV preconditioning is competitive with respect to the Block Jacobi. However the sparsity of the preconditioned system is partially worsened by the low value of the drop tolerance imposed. Either by increasing the value of ε or by applying a filtering on the computed preconditioner, the condition $\rho(H^*)$ is lost as concerns the SP_1 matrices available.

We still aim at applying the AINV preconditioner on other sets of problems, where the properties of the matrix may facilitate the AINV algorithms in providing us with a preconditioner which guarantees convergence for the Monte Carlo linear solvers.

The ILU preconditioner has been computed as well. However even for ILU(0) the sparsity pattern of the preconditioner is inevitably lost. Therefore it was numerically appealing to keep on studying this kind of preconditioners for the problems considered in this section.

As already said about the AINV preconditioners, we preserve our interest for other categories of problems for which the computation of an ILU preconditioner may actually be profitable from the convergence viewpoint.

Chapter 5

Applicable sets for MC solvers

As we have already seen in Chapter 2, sufficient conditions for the convergence of the Monte Carlo linear are very unlikely to be applicable because of the restricted cases of interest (see [HMY13]). Moreover the necessary and sufficient condition requires the computation of $\rho(H)$ and $\rho(H^*)$. The former may be always guaranteed, once a proper preconditioner is picked (e.g. Approximate Inverse). The latter is more problematic to be verified. Therefore in most of the cases it should be checked every time.

The computational burden associated with the spectral radius of H^* has the same complexity as a deterministic algorithm employed to compute the solution to the linear system 1.1 we are interested in.

In order to bypass this task, it is necessary to look for set of matrices for which the Monte Carlo linear solver always converges, maybe with the use of an ad-hoc preconditioner.

5.1 Strictly diagonally dominant matrices

One of the types of matrices for which the MC solver always converges is represented by strictly diagonally dominant matrices.

Definition 1 A matrix $A \in \mathbb{R}^{n \times n}$ is strictly diagonally dominant (s.d.d) by rows if

$$|a_{ij}| > \sum_{\substack{i=1 \\ i \neq j}}^{i=n} |a_{ij}| \quad (5.1)$$

Definition 2 A matrix $A \in \mathbb{R}^{n \times n}$ is strictly diagonally dominant (s.d.d) by columns if

$$|a_{ij}| > \sum_{\substack{j=1 \\ j \neq i}}^{j=n} |a_{ij}| \quad (5.2)$$

When 5.1 holds we can resort to a left diagonal preconditioning and we get an iteration matrix $H = I - D^{-1}A$ such that $\|H\|_\infty < 1$.

Introducing a MAO transition probability for the Forward Method we get

$$P_{ij} = \frac{|H_{ij}|}{\sum_{k=1}^n |H_{ik}|}$$

therefore the entries of H^* are defined as follows

$$H_{ij}^* = \frac{H_{ij}^2}{P_{ij}} = |H_{ij}| \left(\sum_{k=1}^n |H_{ik}| \right).$$

Therefore

$$\sum_{j=1}^n |H_{ij}^*| = \sum_{j=1}^n H_{ij}^* = \left(\sum_{j=1}^n |H_{ij}| \right) \left(\sum_{k=1}^n |H_{ik}| \right) = \left(\sum_{j=1}^n |H_{ij}| \right)^2 < 1 \quad \forall i = 1, \dots, n.$$

This automatically implies that $\rho(H^*) \leq \|H^*\|_\infty < 1$. Thus the Forward Monte Carlo converges. However nothing a priori can be stated about the convergence of the Adjoint method.

Instead if 5.2 holds we can resort to a right diagonal preconditioning and we get an iteration matrix $H = I - AD^{-1}$ such that $\|H\|_1 < 1$. In this case, by following a similar reasoning to the one made before, we conclude that the Adjoint Method converges because $\|H^*\|_1 < 1$, whereas nothing can be said in advance as concerns the Forward method.

5.1.1 Numerical results

In this section of numerical experiments we focus on a set of matrices obtained by a diagonal shift of the matrices associated with SP_1 , SP_3 and SP_5 equations. The shift applied is such that

$$(A + sD), \quad s \in \mathbb{R}^+, \quad D = \text{diag}(A).$$

All the matrices have been turned into strictly diagonally dominant by columns. Therefore the Adjoint Monte Carlo is expected to converge in all the cases.

In Table 5.1 we report the values of the spectra of H and H^* . For the sake of completeness we have computed the values for both the Forward and the Adjoint Monte Carlo but we will focus on the latter for the computations.

The setting adopted for the parameters of the Adjoint MCSA algorithm is as follows:

- residual relative tolerance: $\varepsilon_1 = 10^{-7}$
- almost optimal transition probability
- maximal # steps per history: 10
- statistical error-based adaptive parameter: $\varepsilon_2 = 0.5$
- granularity of the adaptive approach: $n_{histories} = 1000$

Numerical results are shown in 5.2.

Initial matrix	s	$\rho(H)$	$Forward - \rho(H^*)$	$Adjoint - \rho(H^*)$
SP_1 (a)	0.3	0.7597	0.7441	0.6983
SP_1 (b)	0.4	0.7046	1.1448	0.5680
SP_3	0.9	0.5869	0.4426	0.3727
SP_5	1.6	0.5477	0.3790	0.3431

Table 5.1: S.d.d. diagonally shifted SP_N . Spectra of H and H^* - Forward and Adjoint method

matrix	nnz	s	relative err.	# iterations	CPU time (s)
SP_1 (a)	486438	0.3	$6.277 \cdot 10^{-7}$	36	2271 ($\approx 40\text{min}$)
SP_1 (b)	998631	0.4	$9.885 \cdot 10^{-7}$	21	1172 ($\approx 15\text{min}$)
SP_3	846549	0.9	$5.919 \cdot 10^{-7}$	18	3663 ($\approx 40\text{min}$)
SP_5	1399134	1.6	$4.031 \cdot 10^{-7}$	19	6491 ($\approx 110\text{min}$)

Table 5.2: S.d.d. diagonally shifted SP_N . Adjoint MCSA. Numerical results

The condition of strictly diagonally dominance is a sufficient condition to gain the converge of the algorithm. Nevertheless it is not necessary at all. Indeed, by reducing the entity of the diagonal shift, we lose the s.d.d. property but the stochastic linear solver provides an accurate solution as well. In Table 5.3 and 5.4 we have respectively the spectral radii and the numerical results. The spectral radii have increased with respect to the previous set of cases. In fact the property of diagonally dominance sped up the computations. Once this property is missing, the convergence of the algorithm may be preserved but the number of numerical iterations and the total time employed upsurge.

Initial matrix	s	$\rho(H)$	$Forward - \rho(H^*)$	$Adjoint - \rho(H^*)$
SP_1 (a)	0.2	0.8230	0.8733	0.8195
SP_1 (b)	0.2	0.8220	1.5582	0.7731
SP_3	0.3	0.8126	0.9459	0.7961
SP_5	0.7	0.8376	0.8865	0.8026

Table 5.3: Diagonally shifted SP_N . Spectra of H and H^* - Forward and Adjoint method

matrix	s	relative err.	# iterations	CPU time (s)
SP_1 (a)	0.2	$6.394 \cdot 10^{-7}$	48	3438 ($\approx 1h$)
SP_1 (b)	0.2	$2.59 \cdot 10^{-6}$	36	1581 ($\approx 30min$)
SP_3	0.3	$5.35 \cdot 10^{-7}$	45	9715 ($\approx 150min$)
SP_5	0.7	$4.21 \cdot 10^{-7}$	70	22696 ($\approx 6h$)

Table 5.4: Diagonally shifted SP_N . Adjoint MCSA. Numerical results

5.2 Generalized diagonally dominant matrices (GDDM)

Another set of matrices for which it is possible to guarantee convergence of the Monte Carlo linear solvers is represented by *generalized diagonally dominant matrices*.

Definition 3 A square matrix $A \in \mathbb{R}^{n \times n}$ is said to be *generalized diagonally dominant* if

$$|a_{ii}|x_i \geq \sum_{\substack{j=1 \\ j \neq i}}^n |a_{ij}|x_j, \quad i = 1, \dots, n$$

for some positive vector $\mathbf{x} = (x_1, \dots, x_n)^T$.

A proper subset of the generalized diagonally dominant matrices is represented by M -matrices (see [Axe96]).

There are many definitions of M -matrices that may be found on books and all of them are equivalent. The one presented in [Saa03] states the following

Definition 4 A matrix $A \in \mathbb{R}^{n \times n}$ is said to be an M -matrix if it satisfies the following four properties:

- $a_{i,i} > 0$ for $i = 1, \dots, n$
- $a_{i,j} \leq 0$ for $i \neq j$, $i, j = 1, \dots, n$
- A is nonsingular
- $A^{-1} \geq 0$

The following theorem is useful in order to build a convergent splitting.

Theorem 6 Let $A \in \mathbb{R}^{n \times n}$ be a M -matrix that is partitioned in block matrix form. Then

- The matrices $A_{i,i}$ on the diagonal of A are M -matrices
- The block lower and upper triangular part of A are M -matrices

It implies that Jacobi, Block Jacobi and Gauss Seidel are convergent splitting for an M -matrix. However, all this useful properties are not enough to ensure the convergence of the Monte Carlo linear solver. In fact none of these conditions is sufficient to guarantee $\rho(H^*) < 1$.

In order to restore convergence for a generalized diagonally dominant matrix we used a result shown in [Li02]. In this paper the author present an algorithm to transform a generalized diagonally dominant matrix with nonzero diagonal entries into strictly diagonally dominant. The algorithm works for a generic complex matrix $A \in \mathbb{C}^{n \times n}$.

Here below we report the algorithm at hand

For a given complex matrix A , $a_{ii} \neq 0$, $i = 1, \dots, n$;

1. Compute $S_i = \sum_{j=1, j \neq i}^n |a_{ij}|$, $i = 1, 2, \dots, n$
2. Set $t = 0$. For $i = 1, 2, \dots, n$, if $|a_{ii}| > S_i$, then set $t = t + 1$
3. If $t = 0$, then print "A is not a GDDM": END
4. If $t = n$, then print "A is a GDDM": END
5. **for** $i=1, n$ **do**

$$d_i = \frac{S_i + \varepsilon}{|a_{ii}| + \varepsilon} \quad \varepsilon > 0, \quad j = 1, 2, \dots, n;$$

$$a_{ji} = a_{ji} \cdot d_i$$
end
6. Go to step 1.

Algorithm 7: Algorithm to turn a GDDM matrix into a sdd by rows.

This approach turns a generalized diagonally dominant matrix into a strictly diagonally dominant matrix by rows. By substituting $S_i = \sum_{j=1, j \neq i}^n |a_{ij}|$ at step 1 with $S_i = \sum_{i=1, i \neq j}^n |a_{ij}|$ and by replacing $a_{ji} = a_{ji} \cdot d_i$ with $a_{ji} = a_{ji} \cdot d_j$ we obtain the algorithm that turns a GDDM matrix into a s.d.d. by columns.

Once we have applied this transformation to the matrix at hand into a s.d.d., we can use the Monte Carlo linear solver which is ensured to converge.

5.3 "Block-diagonally dominant"-like matrices

In Chapter 4 we have already discussed the Block diagonal preconditioning. In this section we wonder in which situations a preconditioning such as this guarantees to come up with a convergent Monte Carlo linear solver. By mimicking the computations we already showed previously, the iteration matrix $H \in \mathbb{R}^{n \times n}$ resulting from a block diagonal preconditioning assumes the form

$$H = I - D^{-1}A = \begin{bmatrix} 0_{n_1 \times n_1} & -A_{11}^{-1}A_{12} & \cdots & \cdots & -A_{11}^{-1}A_{1p} \\ -A_{22}^{-1}A_{21} & 0_{n_2 \times n_2} & -A_{22}^{-1}A_{23} & \cdots & -A_{22}^{-1}A_{2p} \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ -A_{pp}^{-1}A_{p1} & \cdots & \cdots & -A_{pp}^{-1}A_{p(p-1)} & 0_{n_p \times n_p} \end{bmatrix}.$$

5.3.1 Forward method

By assuming that all n_i have the same value we may define

$$m = \text{size of a block} = \frac{n_i}{p}.$$

The MAO transition probability matrix assumes the following form:

$$P_{i,j} = \frac{|H_{i,j}|}{\sum_{k=1}^n |H_{i,k}|} = \frac{\left| \left(A_{\lfloor \frac{i}{m} \rfloor \lfloor \frac{i}{m} \rfloor}^{-1} A_{\lfloor \frac{i}{m} \rfloor \lfloor \frac{j}{m} \rfloor}^{-1} \right)_{(i \% m)(j \% m)} \right|}{\sum_{k=1, k \neq i}^n \left| \left(A_{\lfloor \frac{i}{m} \rfloor \lfloor \frac{i}{m} \rfloor}^{-1} A_{\lfloor \frac{i}{m} \rfloor \lfloor \frac{k}{m} \rfloor}^{-1} \right)_{(i \% m)(k \% m)} \right|}$$

Consequently, the H^* matrix is defined such that

$$H_{i,j}^* = |H_{i,j}| \left(\sum_{k=1}^n |H_{ik}| \right) = \left| \left(A_{\lfloor \frac{i}{m} \rfloor \lfloor \frac{i}{m} \rfloor}^{-1} A_{\lfloor \frac{i}{m} \rfloor \lfloor \frac{j}{m} \rfloor}^{-1} \right)_{(i \% m)(j \% m)} \right| \left| \sum_{\substack{k=1 \\ k \neq i}}^n \left| \left(A_{\lfloor \frac{i}{m} \rfloor \lfloor \frac{i}{m} \rfloor}^{-1} A_{\lfloor \frac{i}{m} \rfloor \lfloor \frac{k}{m} \rfloor}^{-1} \right)_{(i \% m)(k \% m)} \right| \right|$$

By computing the sum over a generic row of H^* we get:

$$\sum_{j=1}^n |H_{i,j}^*| = \sum_{j=1}^n H_{i,j}^* = \left(\sum_{\substack{j=1 \\ j \neq i}}^n \left| \left(A_{\lfloor \frac{i}{m} \rfloor \lfloor \frac{i}{m} \rfloor}^{-1} A_{\lfloor \frac{i}{m} \rfloor \lfloor \frac{j}{m} \rfloor}^{-1} \right)_{(i \% m)(j \% m)} \right| \right)^2$$

if we focus on the norm $\|H^*\|_\infty$, then the following equivalence condition holds:

$$\|H^*\|_\infty < 1 \Leftrightarrow \sum_{\substack{j=1 \\ j \neq i}}^n \left| \left(A_{\lfloor \frac{i}{m} \rfloor \lfloor \frac{i}{m} \rfloor}^{-1} A_{\lfloor \frac{i}{m} \rfloor \lfloor \frac{j}{m} \rfloor}^{-1} \right)_{(i \% m)(j \% m)} \right| < 1 \quad \forall i = 1, \dots, n$$

A sufficient condition for this to happen is

$$\sum_{\substack{j=1 \\ j \neq i}}^p \|A_{ii}^{-1} A_{ij}\|_\infty < 1. \quad (5.3)$$

By defining a matrix $\tilde{H} \in \mathbb{R}^{p \times p}$ such that

$$\tilde{H} = \begin{bmatrix} 0 & \|A_{11}^{-1} A_{12}\|_\infty & \cdots & \cdots & \|A_{11}^{-1} A_{1p}\|_\infty \\ \|A_{22}^{-1} A_{21}\|_\infty & 0 & \|A_{22}^{-1} A_{23}\|_\infty & \cdots & \|A_{22}^{-1} A_{2p}\|_\infty \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \|A_{pp}^{-1} A_{p1}\|_\infty & \cdots & \cdots & \|A_{pp}^{-1} A_{p(p-1)}\|_\infty & 0 \end{bmatrix}$$

We can use the \tilde{H} matrix just defined in order to introduce a sufficient condition for the convergence of the Forward Monet Carlo method with a Block Diagonal preconditioning.

$$\|\tilde{H}\|_\infty < 1 \Rightarrow \|H^*\|_\infty < 1. \quad (5.4)$$

5.3.2 Adjoint method

Analogously to the Forward method, if we define

$$(H^*)_{i,j}^T = |H_{i,j}^T| \left(\sum_{k=1}^n |H_{ik}^T| \right)$$

we can formulate a sufficient condition for the convergence of the Adjoint Monte Carlo method by introducing a matrix \tilde{H} which in this case is such that

$$\tilde{H} = \begin{bmatrix} 0 & \|A_{11}^{-1} A_{12}\|_1 & \cdots & \cdots & \|A_{11}^{-1} A_{1p}\|_1 \\ \|A_{22}^{-1} A_{21}\|_1 & 0 & \|A_{22}^{-1} A_{23}\|_1 & \cdots & \|A_{22}^{-1} A_{2p}\|_1 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \|A_{pp}^{-1} A_{p1}\|_1 & \cdots & \cdots & \|A_{pp}^{-1} A_{p(p-1)}\|_1 & 0 \end{bmatrix}.$$

The sufficient condition assume the form

$$\|\tilde{H}\|_1 < 1 \Rightarrow \|H^*\|_1 < 1. \quad (5.5)$$

Bibliography

- [AAD⁺05] V. Alessandrov, E. Atanassov, I. Dimov, S. Branford, A. Thandavan, and C. Weihrauch. Parallel hybrid Monte Carlo algorithms for matrix computations problems. *Lecture Notes in Computer Science*, 3516:752–759, 2005.
- [Axe96] O. Axelsson. *Iterative Solution Methods*. Cambridge University Press, 1996.
- [Ben02] M. Benzi. Preconditioning Techniques for Large Linear Systems: A Survey. *Journal of Computational Physics*, 182:418–477, 2002.
- [BMM96] M. Benzi, C.D. Meyer, and Tuma M. A sparse approximate inverse preconditioner for the conjugate gradient method. *SIAM Journal on Scientific Computing*, 17:1135–1149, 1996.
- [BT98] M. Benzi and M. Tuma. A sparse approximate inverse preconditioner for nonsymmetric linear systems. *SIAM Journal on Scientific Computing*, 19:968–994, 1998.
- [DA98] I.T. Dimov and V.N. Alexandrov. A new highly convergent Monte Carlo method for matrix computations. *Mathematics and Computers in Simulation*, (47):165–181, 1998.
- [DAK01] I. Dimov, V. Alexandrov, and A. Karaivanova. Parallel resolvent Monte Carlo algorithms for linear algebra problems. *Mathematics and Computers in Simulation*, 55(55):25–35, 2001.
- [EMSH14] T.M Evans, S.W. Mosher, S.R. Slattery, and S.P. Hamilton. A Monte Carlo synthetic-acceleration method for solving the thermal radiation diffusion equation. *Journal of Computational Physics*, 258(November 2013):338–358, 2014.
- [ESW13] T.M. Evans, S.R. Slattery, and P.P.H Wilson. A spectral analysis of the domain decomposed Monte Carlo method for linear systems. *International Conference on Mathematics and Computational Methods Applied to Nuclear Science & Engineering*, 2013.
- [Hal62] J.H. Halton. Sequential Monte Carlo. *Mathematical Proceedings of the Cambridge Philosophical Society*, 58(1):57–58, 1962.
- [Hal94] J.H. Halton. Sequential Monte Carlo techniques for the solution of linear systems. *Journal of Scientific Computing*, 9(2):213–257, 1994.
- [HJ94] R. A. Horn and C. Johnson. *Topics in Matrix Analysis*. Cambridge University Press, 1994.
- [HMY13] J. Hao, M. Mascagni, and L. Yaohang. Convergence analysis of Markov chain Monte Carlo linear solvers using Ulam-von Neumann algorithm. *SIAM Journal on Numerical Analysis*, 51(4):2107–2122, 2013.
- [Li02] L. Li. On the iterative criterion for generalized diagonally dominant matrices. *SIAM Journal in Matrix Analysis and Applications*, 24(1):17–24, 2002.
- [Saa03] Y. Saad. *Iterative Methods for Sparse Linear Systems*. Society for Industrial and Applied Mathematics, 2003.
- [Sla13] S. Slattery. *Parallel Monte Carlo Synthetic Acceleration Methods For Discrete Transport Problems*. PhD thesis, University of Wisconsin-Madison, 2013.